# Introduction to Operating Systems, Part II

**CSD2180 Operating Systems**
BSc in Computer Science (IMGD / RTIS)
Singapore Institute of Technology / DigiPen Institute of Technology
September 2021

# Attendance Taking

**https://forms.office.com/r/t6E6nkGXuA**

- Log in to your SIT account to submit the form
- You can only submit the form once
- The codeword is **durian**

SINGAPORE INSTITUTE OF TECHNOLOGY

# Operating System Operations

# Different Perspectives of the Same Computer Motherboard


**X570S AORUS MASTER Motherboard Layout**

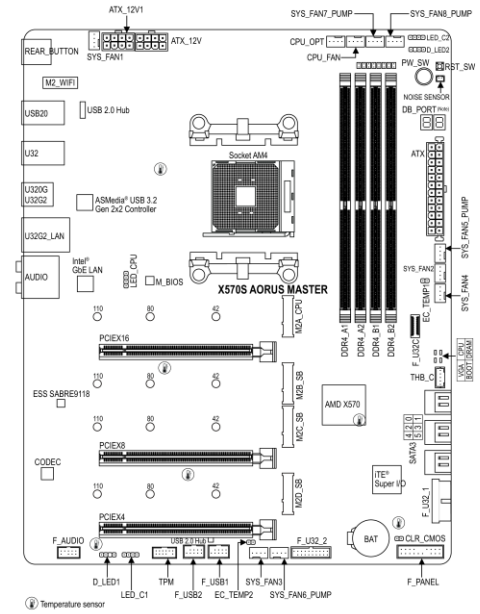


(Note) For debug code information, please refer to Chapter 6.



**X570S AORUS MASTER Motherboard Block Diagram**

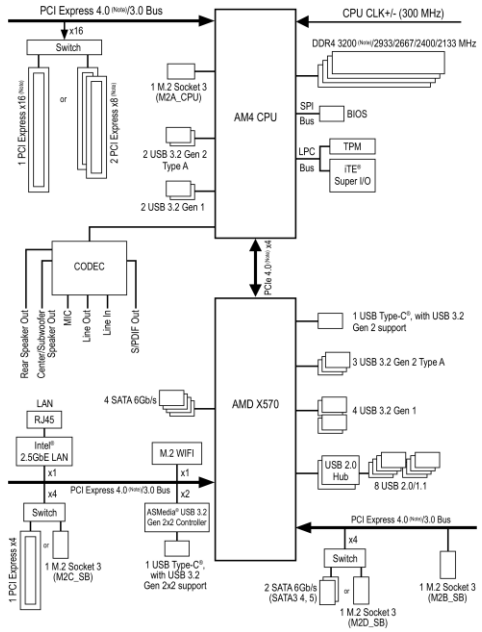


(Note) Actual support may vary by CPU.

# How a Computer Boots Up

## Hardware powered on

Press this button

# How a Computer Boots Up

**Hardware powered on**

**BIOS runs**

- Performs power-on self-test (POST)
- Initialize peripherals
- Locates secondary storage devices for bootable devices

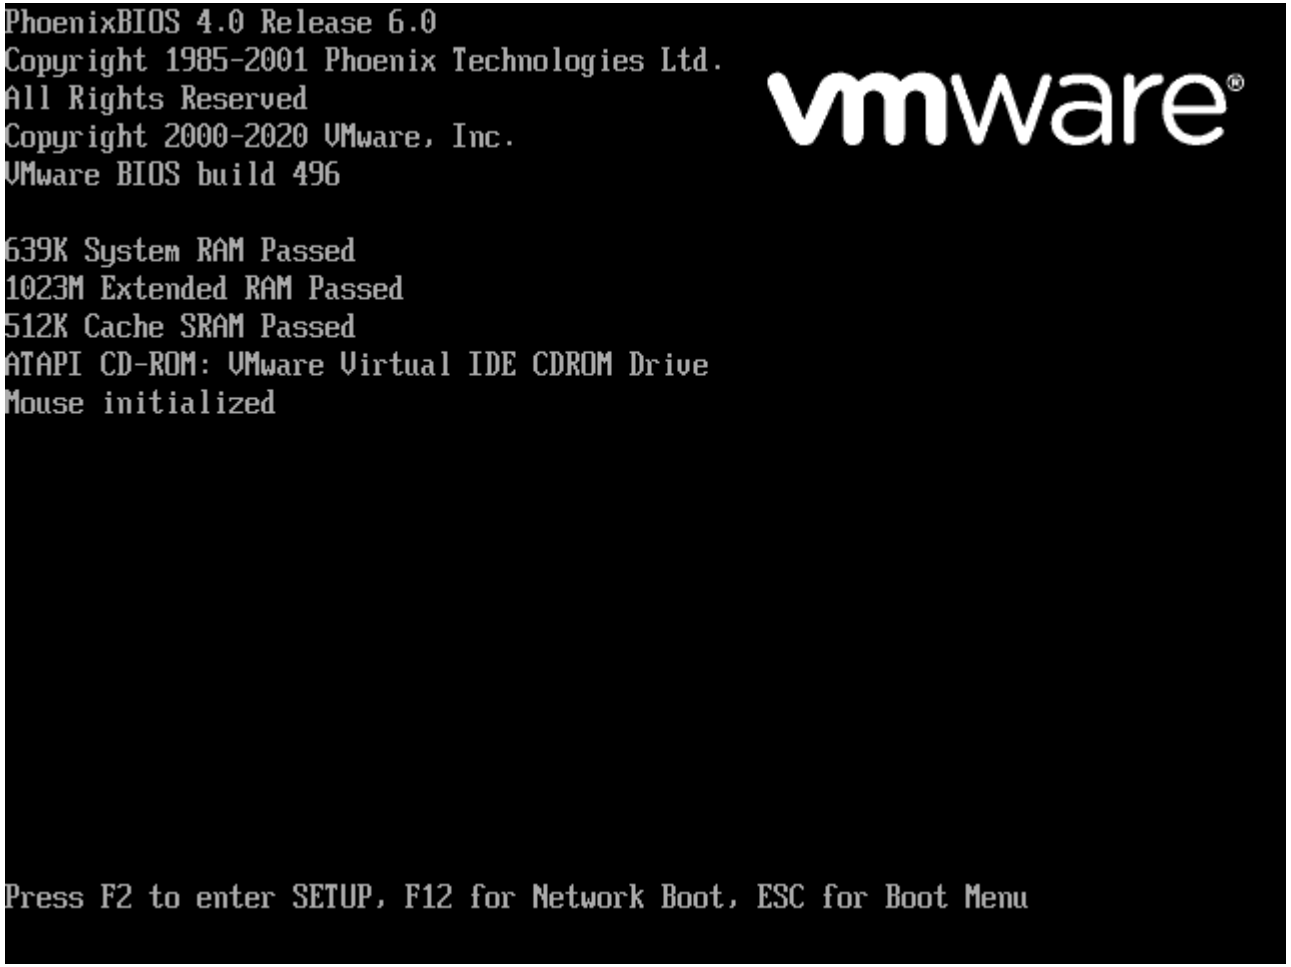

# How a Computer Boots Up

**Hardware powered on**

**BIOS runs**

**BIOS loads and runs Master Boot Record (MBR)**

- Located at first 512 bytes of first / designated boot device

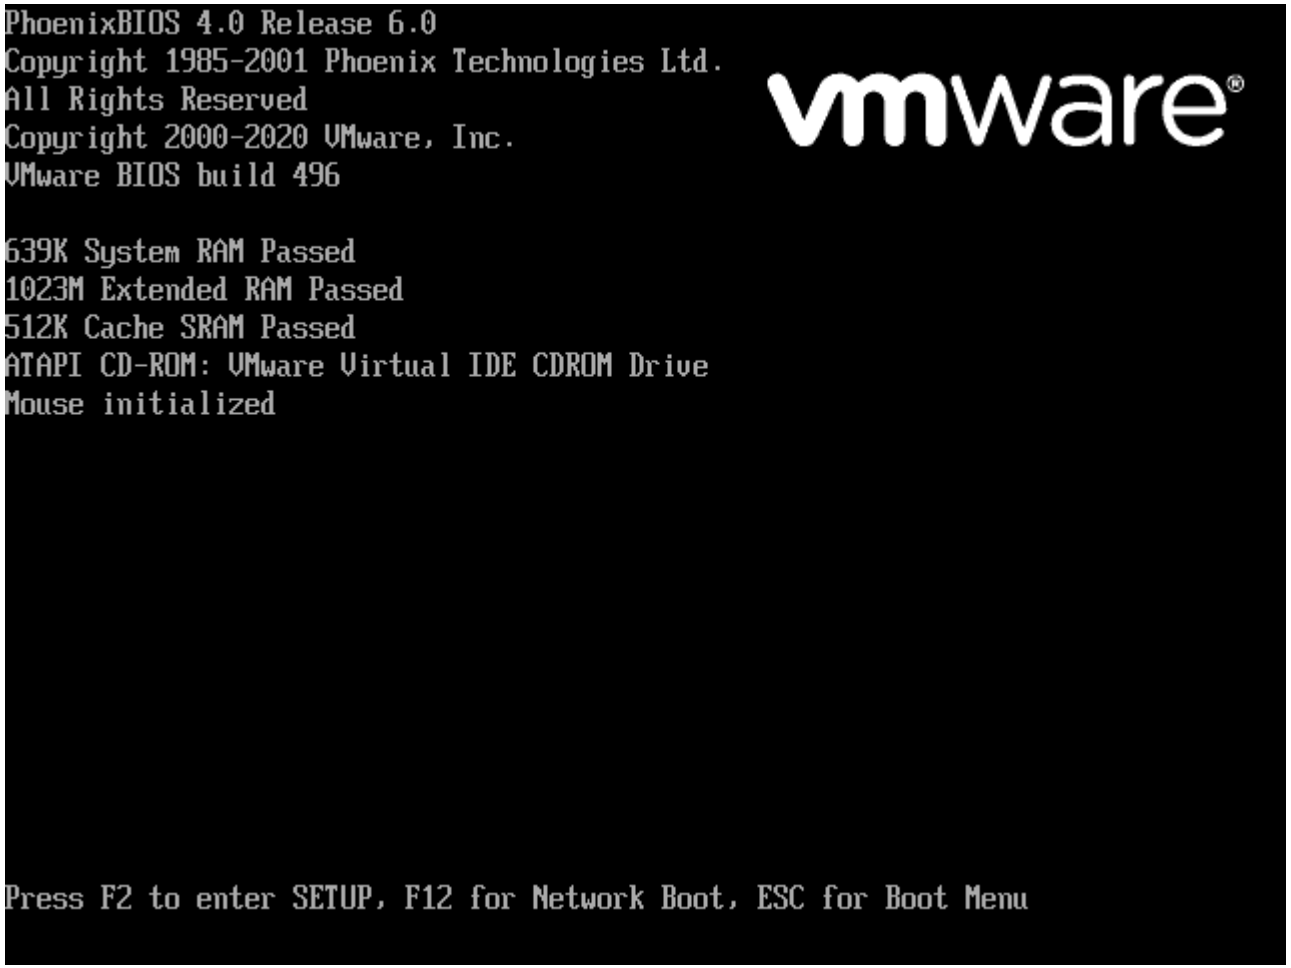

# How a Computer Boots Up

**Hardware powered on**

**BIOS runs**

**BIOS loads and runs Master Boot Record (MBR)**

**MBR loads bootloader**

- Bootloader may be loaded from e.g., partitions
- A bootloader may load subsequent bootloader stages, i.e., chain-loading

# How a Computer Boots Up

- Hardware powered on
- BIOS runs
- BIOS loads and runs Master Boot Record (MBR)
- MBR loads bootloader
- Bootloader loads and runs the operating system

# How a Computer Boots Up

- Hardware powered on
- BIOS runs
- BIOS loads and runs Master Boot Record (MBR)
- MBR loads bootloader
- Bootloader loads and runs the operating system
- Operating system starts up and runs
  - Kernel loads and executes
  - Services are started

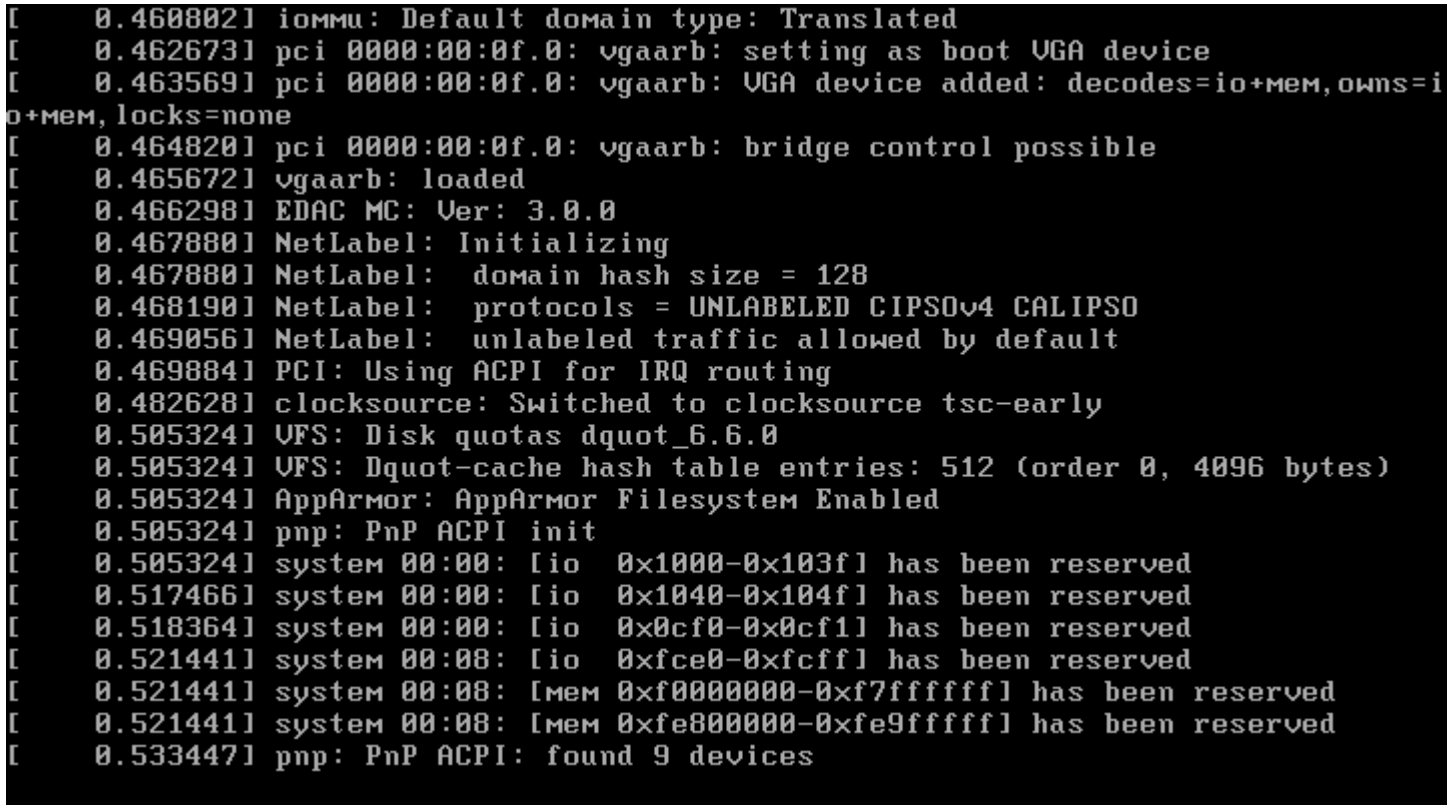

# How a Computer Boots Up

- Hardware powered on
- BIOS runs
- BIOS loads and runs Master Boot Record (MBR)
- MBR loads bootloader
- Bootloader loads and runs the operating system
- Operating system starts up and runs
  - Kernel loads and executes
  - Services are started

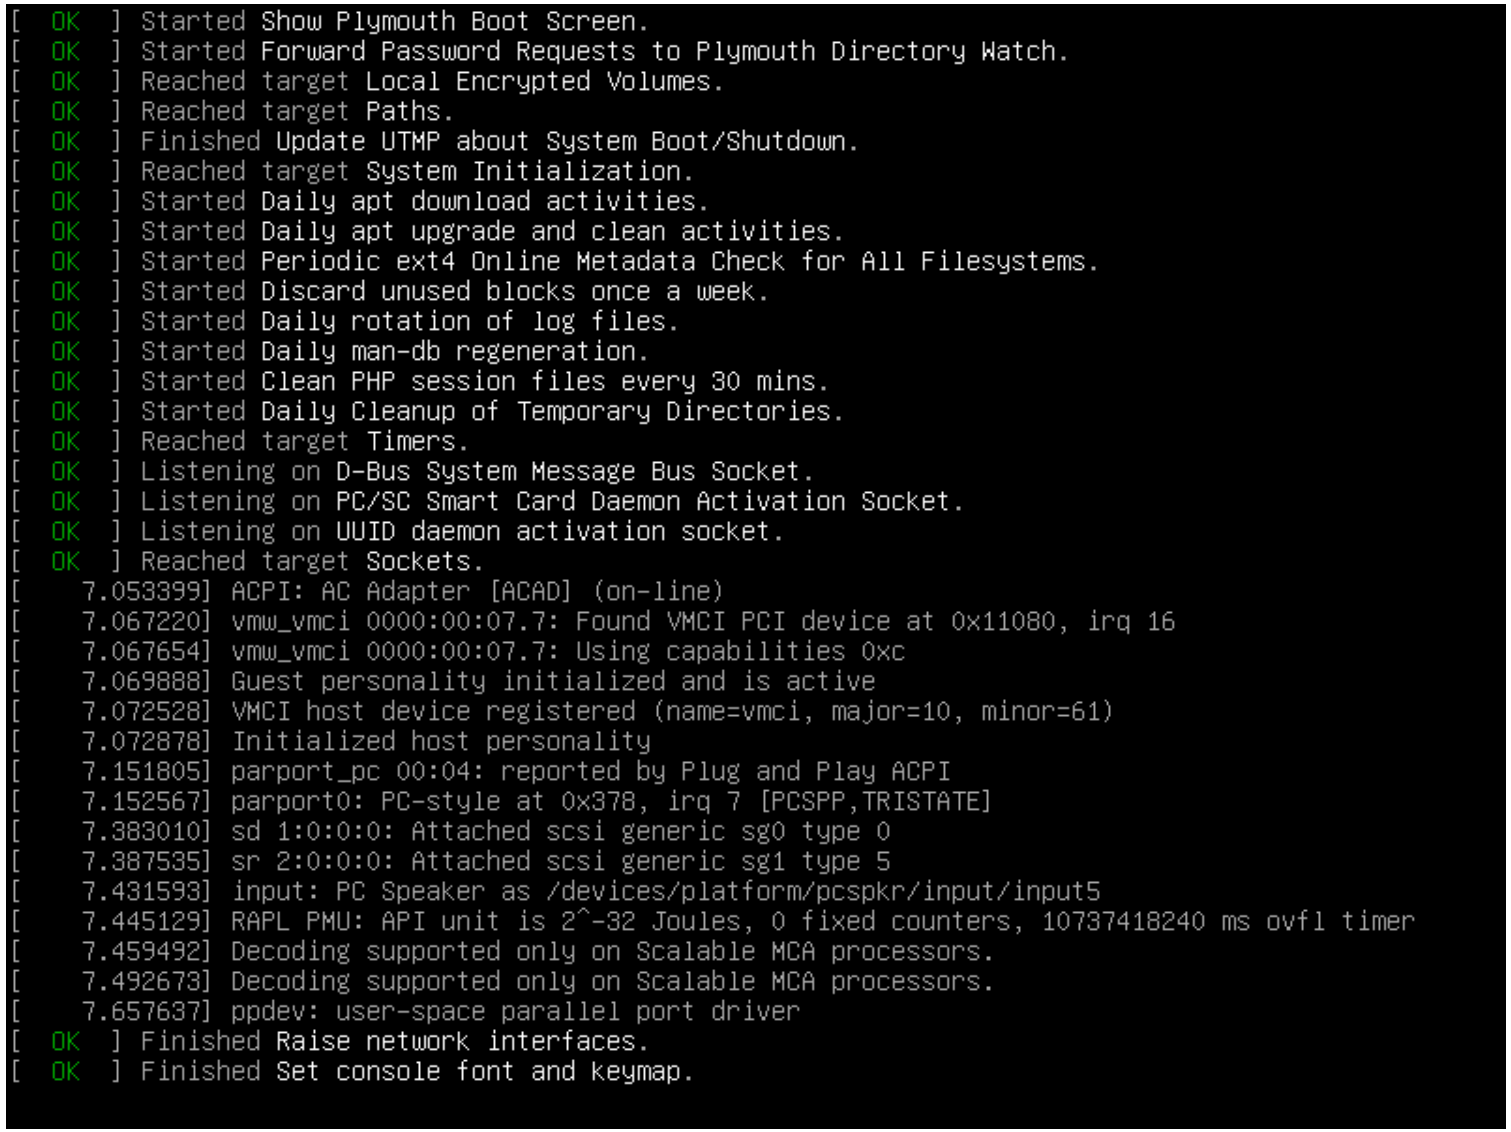

# How a Computer Boots Up

- Hardware powered on
- BIOS runs
- BIOS loads and runs Master Boot Record (MBR)
- MBR loads bootloader
- Bootloader loads and runs the operating system
- Operating system starts up and runs

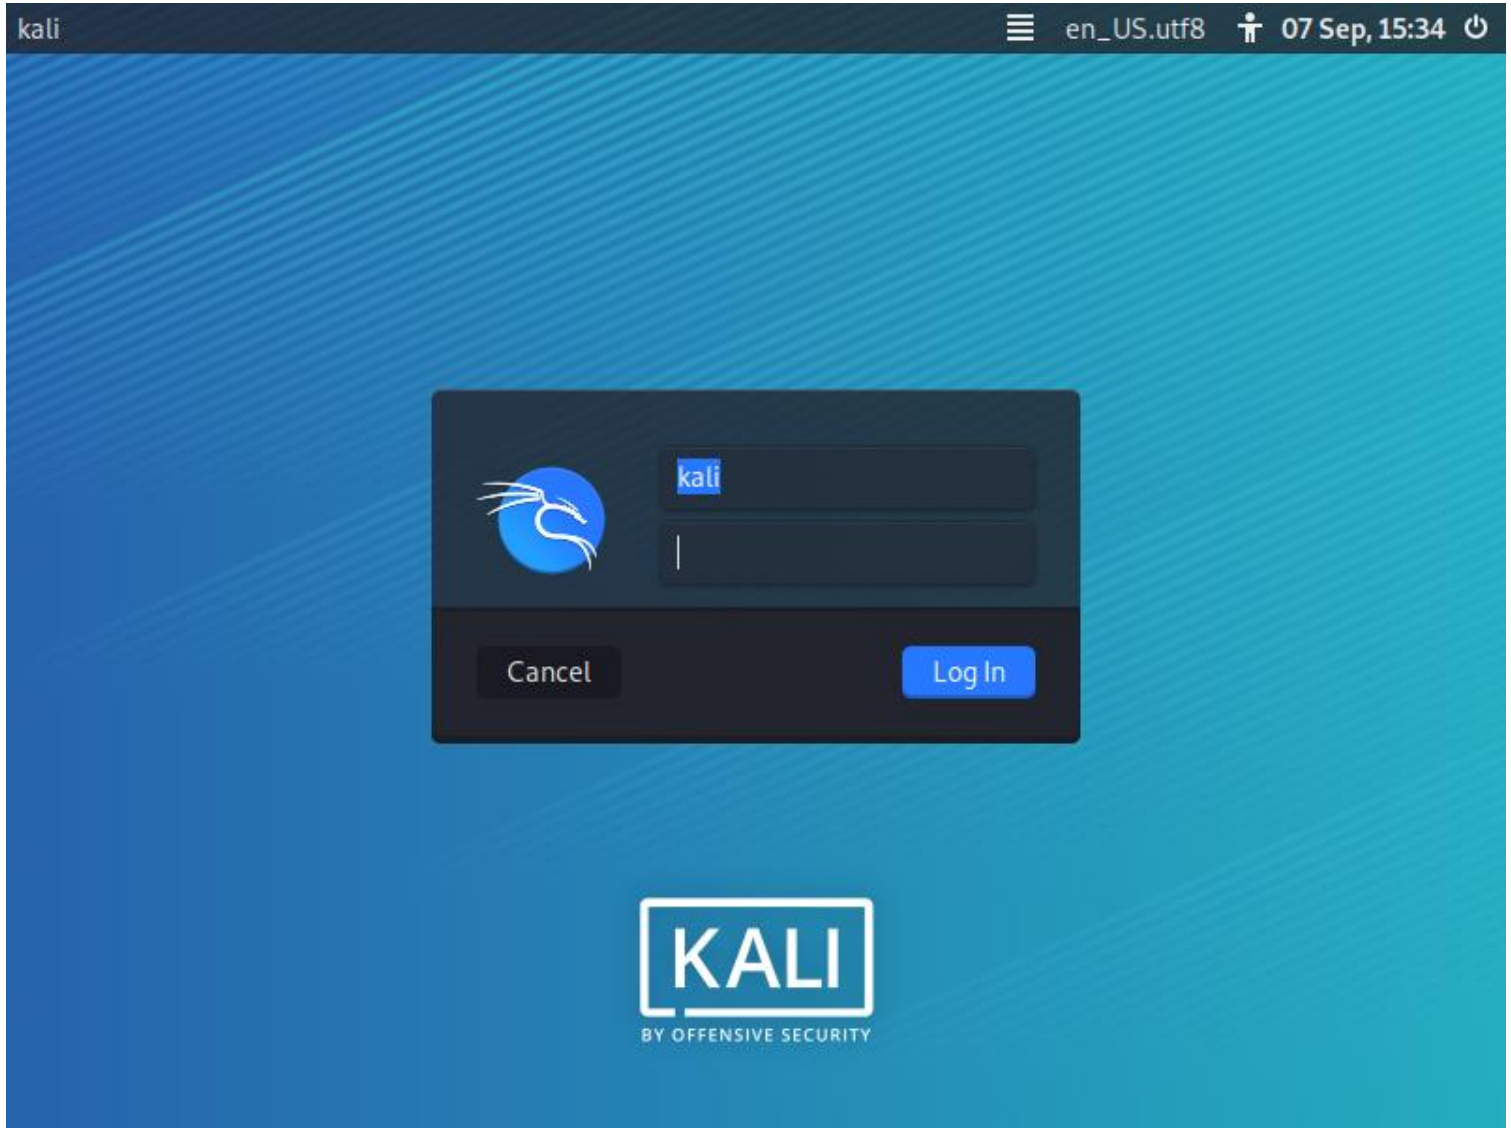

# Multiprogramming and Multitasking

Problem

- A single program cannot always keep the CPU or I/O devices busy
- Users want to run more than one program at a time

**Multiprogramming** is needed for efficiency, by organizing programs so that the CPU always has one to execute

General idea

- The OS keeps several processes in memory
- One process is selected and run via **process scheduling**
- When the process has to wait (e.g., for I/O), the OS switches to another process

**Protip: A *program* in execution is called a *process***

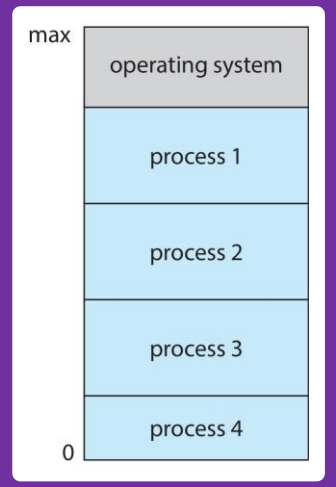

# Multiprogramming and Multitasking

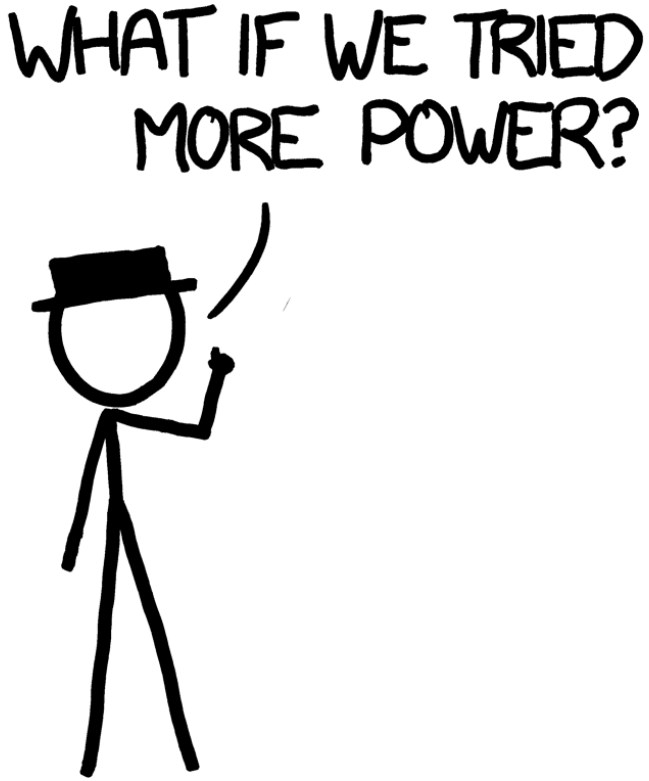

# Multiprogramming and Multitasking

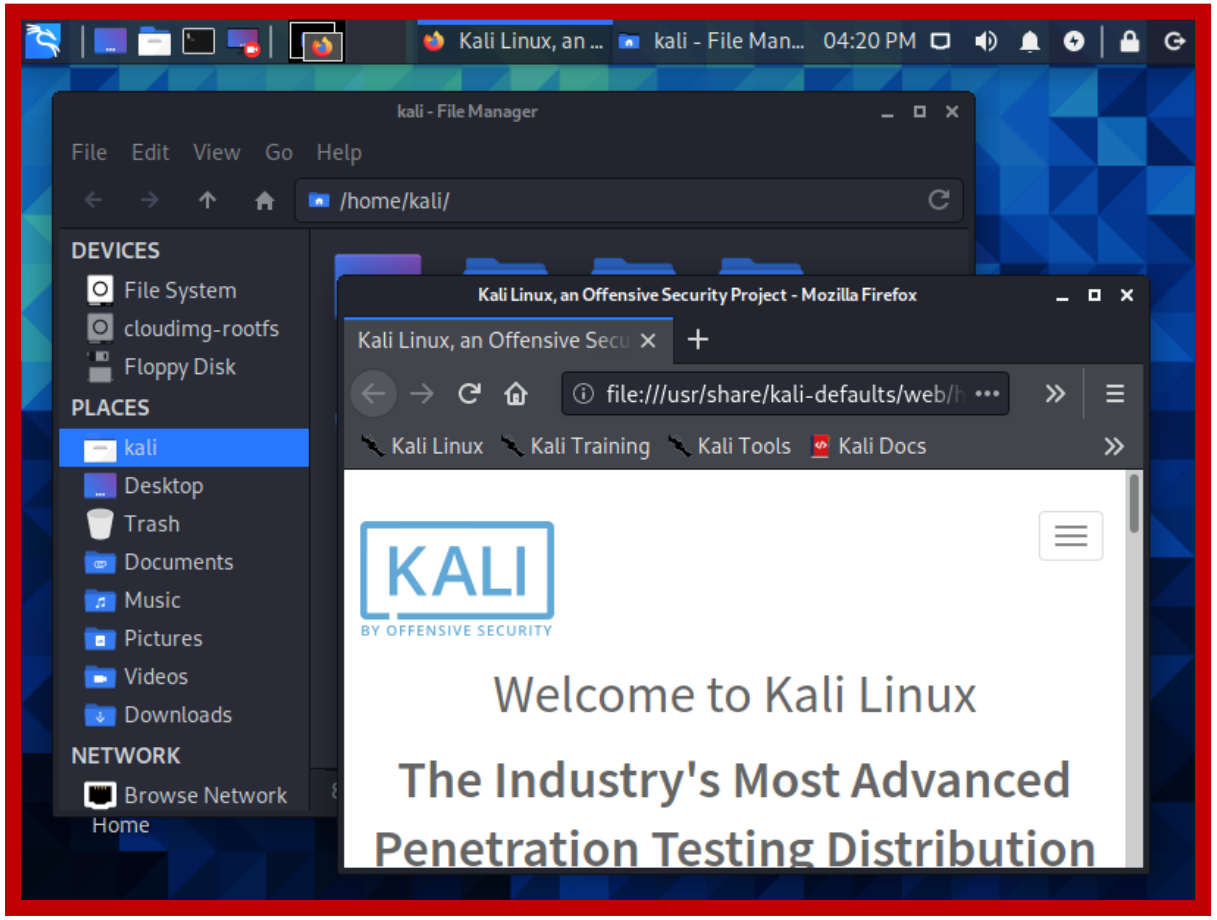


**Multitasking:** The CPU switches jobs so frequently that users can *interact* with each job while it is running, creating **interactive computing** ☺

In general

- Response time should be *fast* (e.g., sub-second)
- User have >1 program executing in memory
- If >1 jobs ready to run at the same time, then **scheduling** takes place
- If processes don't fit in memory, **swapping** moves them in and out to run
- **Virtual memory** allows execution of processes not completely in memory

# Dual-Mode Operation

**User mode** and **kernel mode**

- **Mode bit** provided by hardware
- Provides ability to distinguish when system is running user code or kernel code

Some instructions are designated as **privileged**, only executable in kernel mode

How do we guarantee that user does not explicitly set the mode bit to "kernel"?

- System call sets mode to kernel, and on return from call resets it back to user
- Instruction to switch to kernel mode is handled specially (sometimes via a trap, sometimes as a unique instruction)

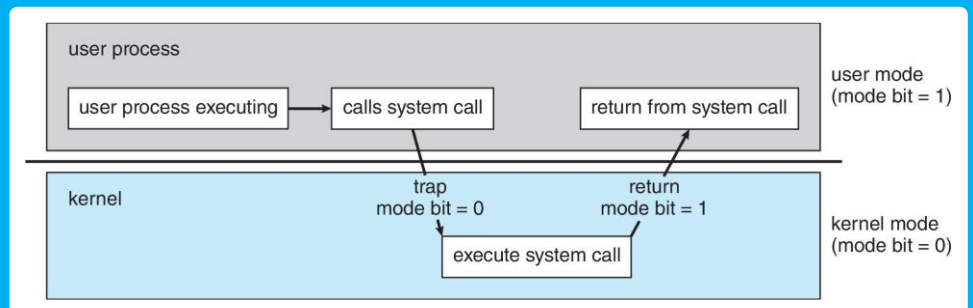


Dual-mode operation allows OS to protect itself and other system components

**Protip: Kernel mode is also called *supervisor mode*, *system mode*, or *privileged mode***

# Timer

**Timer to prevent infinite loop or process hogging resources**

Timer is set to interrupt the computer after some time period

Keeps a counter that is decremented by the physical clock

The operating system set the counter (privileged instruction)

When counter reaches zero, an interrupt is generated

Set up before scheduling process to regain control or terminate program that exceeds allotted time

SINGAPORE INSTITUTE OF TECHNOLOGY

# Resource Management

# Recap

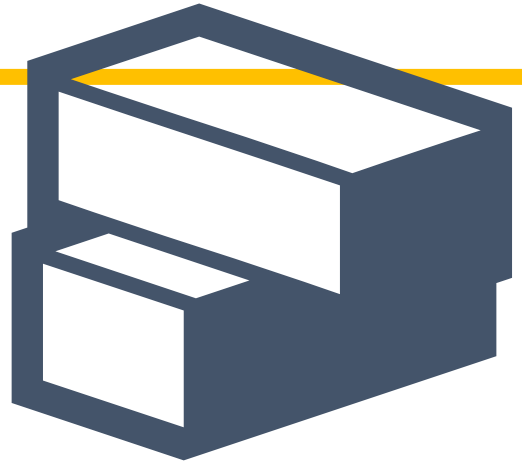

Memory management

Filesystem management

Mass storage management

Caching

***We have discussed some aspects of operating system resource management...***

# Why Are We Talking About Storage?

**Memory Management**

To execute a program all (or part) of the instructions must be in memory

Memory management determines what is in memory and when

Memory management activities

- Keeping track of which parts of memory are currently being used and by whom
- Deciding which processes (or parts thereof) and data to move into and out of memory
- Allocating and deallocating memory space as needed

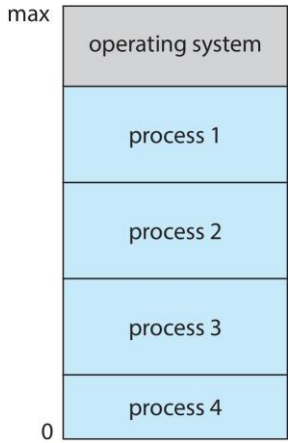

# Why Are We Talking About Storage?

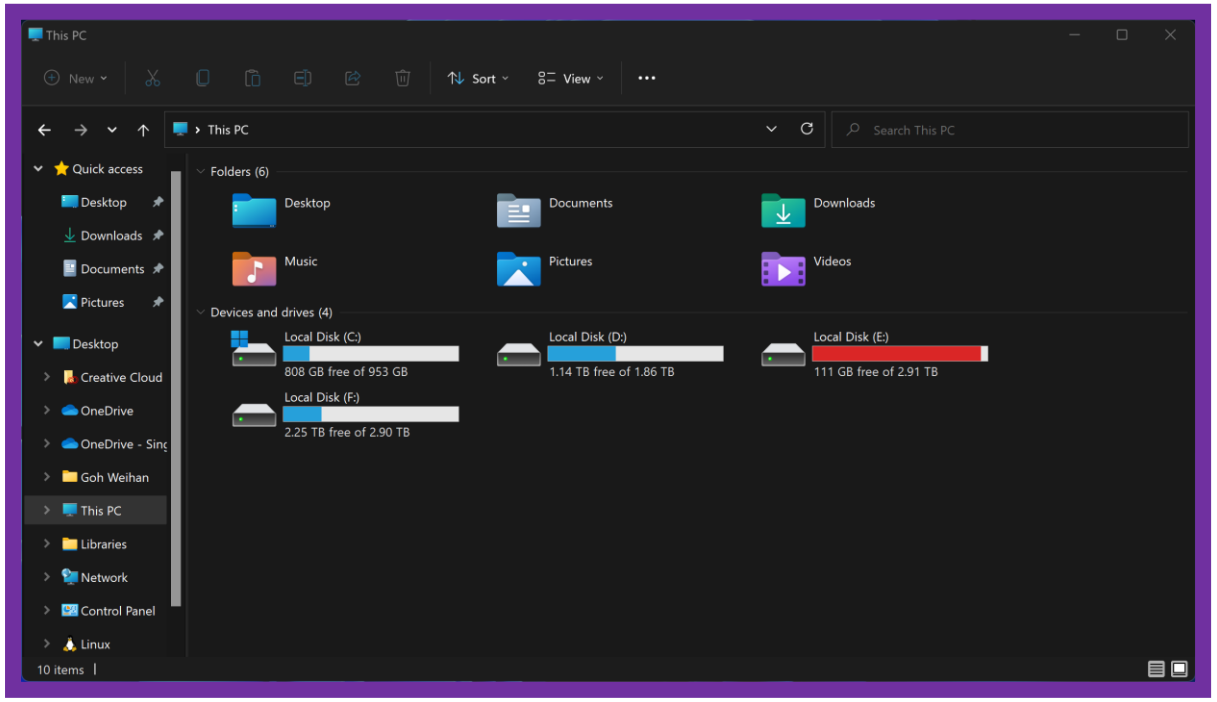


## Filesystem Management

Operating system provides uniform, logical view of information storage

Filesystem management

- Files usually organized into directories
- Access control on most systems to determine who can access what
- Operating system activities include
  - Creating and deleting files and directories
  - Primitives to manipulate files and directories
  - Mapping files onto secondary storage
  - Backup files onto stable (non-volatile) storage media

# Why Are We Talking About Storage?

## Mass-Storage Management

Usually, disks are used to store data that does not fit in main memory, or data that must be kept for a long period of time

Proper management is of central importance

Operating system activities

- Mounting and unmounting (i.e., safe removal)
- Free-space management
- Storage allocation
- Disk scheduling
- Partitioning
- Protection

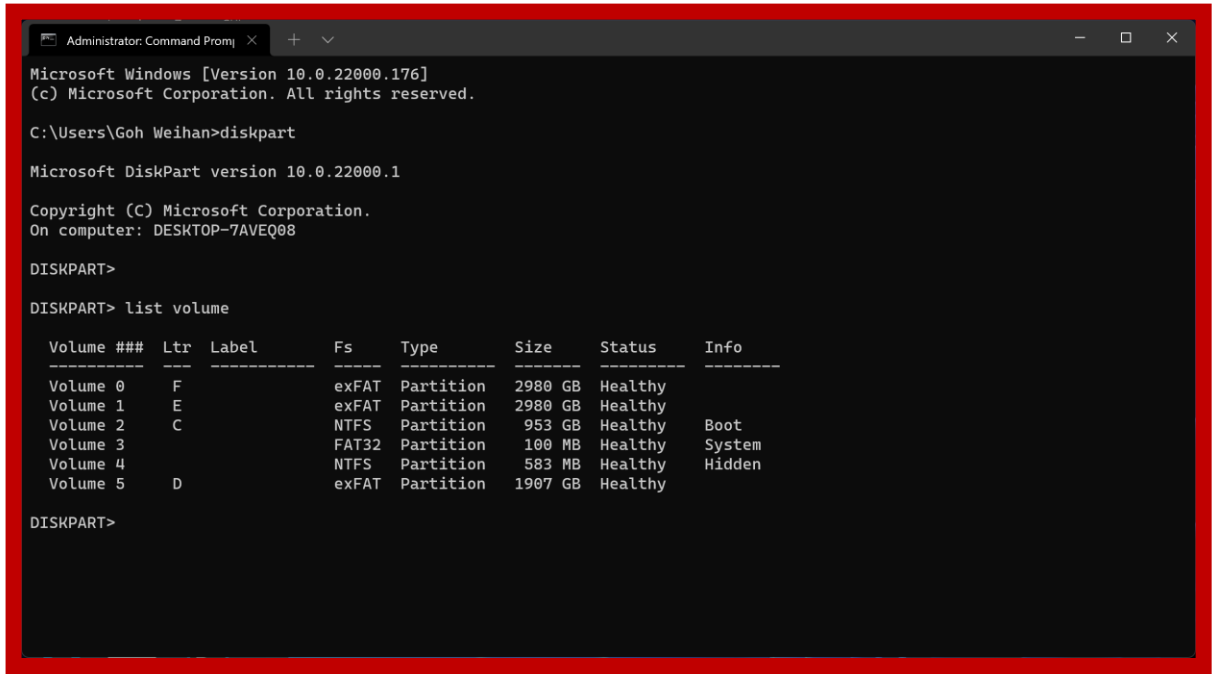

# Why Are We Talking About Storage?

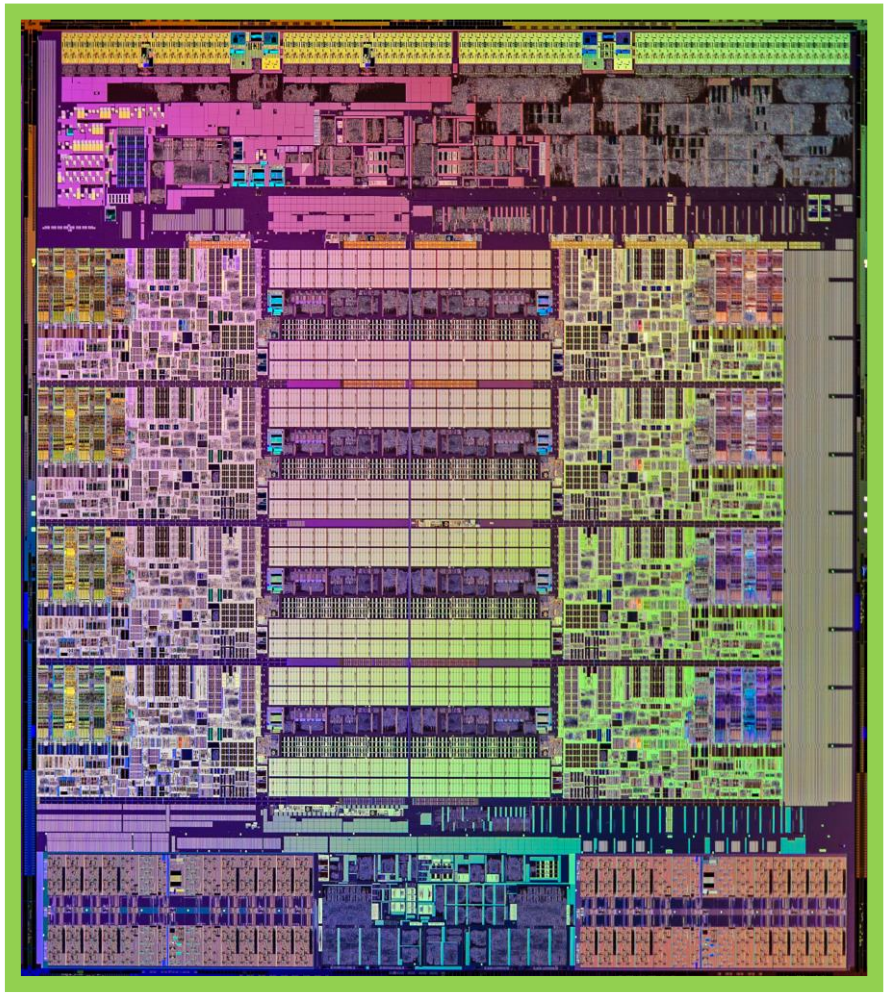

**Caching**

Information in use copied from slower to faster storage temporarily

Faster storage (cache) checked first to determine if information is there

- If it is, information used directly from cache (fast)
- If not, data copied to cache and used there

Cache smaller than storage being cached

- Cache management important design problem
- Cache size and replacement policy

# Process Management

Process needs resources to accomplish its task

- CPU, memory, I/O, files
- Initialization data

Process termination requires reclaim of any reusable resources

Typically, a system has many processes, some user, and the operating system running *concurrently* on one or more CPUs

- Concurrency by multiplexing the CPUs among the processes / threads

**A program in execution is called a process**

- Program = passive entity
- Process = active entity

**Single-threaded process has one *program counter* specifying location of next instruction to execute**

- Process executes instructions "sequentially", one at a time, until completion

**Multi-threaded process has one program counter *per thread***

# Process Management Activities

Create / delete user and system processes

Suspend / resume processes

Provide mechanisms for process synchronization

Provide mechanisms for process communication

Provide mechanisms for deadlock handling

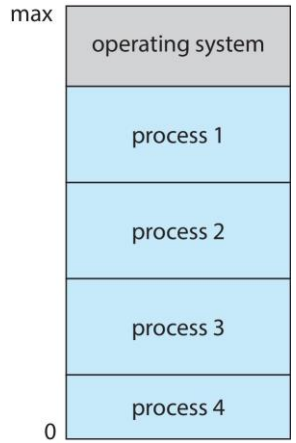

# I/O Subsystem Management

**Hide peculiarities of hardware devices from the user**

**Memory management of I/O, including**

- **Buffering:** Storing data temporarily while it is being transferred
- **Caching:** Storing parts of data in faster storage for performance
- **Spooling:** the overlapping of output of one job with input of other jobs

**General device-driver interface**

**Drivers for specific hardware devices**

SINGAPORE INSTITUTE OF TECHNOLOGY

# Security and Protection

# Security and Protection

**Systems generally distinguish among users to determine who can do what**

- User identities (user IDs, security IDs) include name and associated identifier, one per user
- User identifier then associated with all files and processes of that user to determine access control
- Group identifier (group ID) allows sets of users to be defined and controls managed, then also associated with each process and file

**Privilege escalation allows user to access resources of another user with *more* rights**

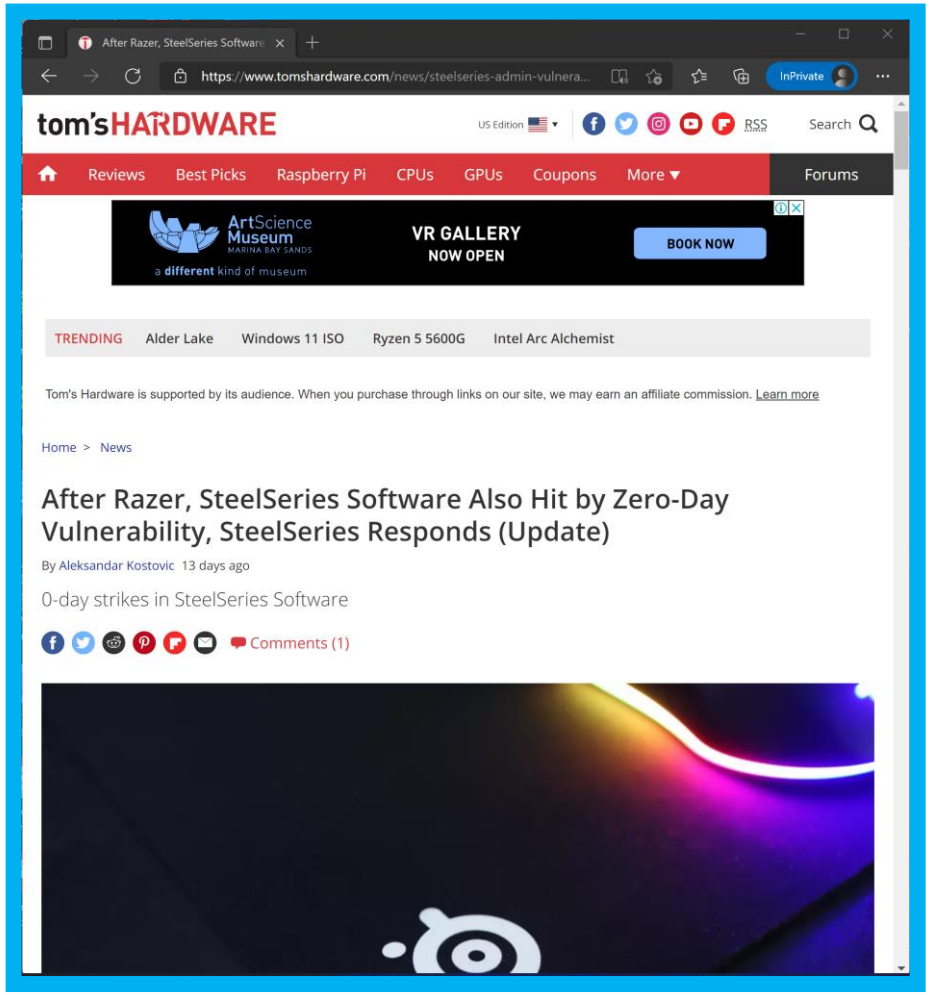

SINGAPORE INSTITUTE OF TECHNOLOGY

# Computing Environments

 Traditional |  Emulated | 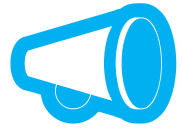
Virtualized | 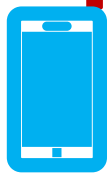 Mobile |  Cloud
Computing |  Real-Time Embedded

# Traditional Computing

Stand-alone, general-purpose machines

Definition kind of blurred these days as most computer systems **interconnect** with others (e.g., over the Internet)

Networking is now ubiquitous

# Emulation

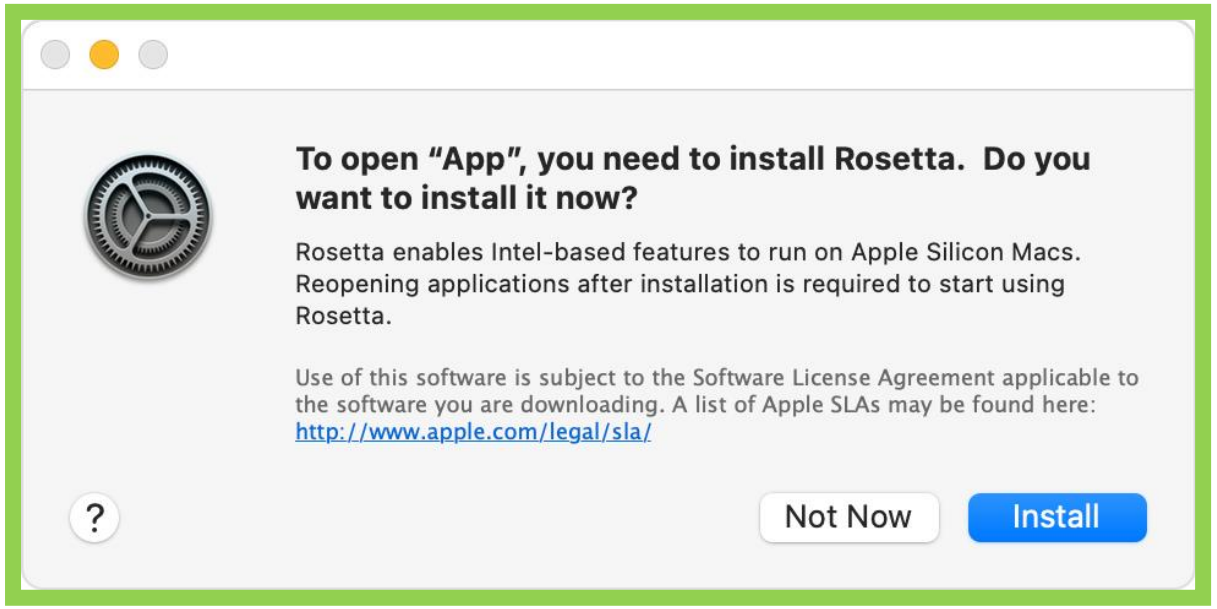


**Emulation** involves simulating computer hardware in software

Typically used for running applications when source CPU type different from target CPU type (e.g., running ARM applications on x86)

Can be (and generally) slow

- Every machine-level instruction that runs natively on the source system must be **translated** to the equivalent on the target system

**Where has this been used before?**

# Virtualization

**Operating system natively compiled for CPU, running guest operating systems that are also natively compiled for CPU**

- For example, using VMware Workstation to run Windows XP guests on a Windows 10 host system
- Virtual machine manager (VMM) provides virtualization services

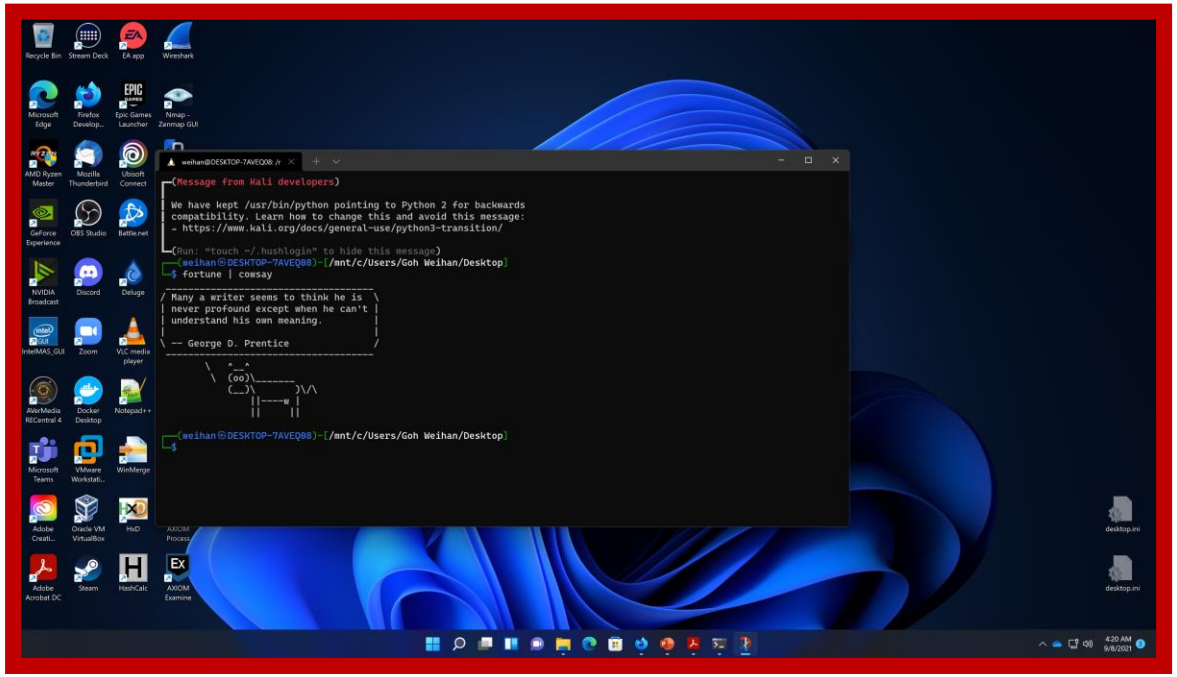

# Mobile Computing

Handheld smartphones, tablets, etc.

**More features**, such as GPS, gyroscope, etc.

Allow new types of **use cases** like augmented reality

Use **wireless or cellular data** networks for connectivity

Ecosystem led by Apple iOS and Google Android

# Cloud Computing

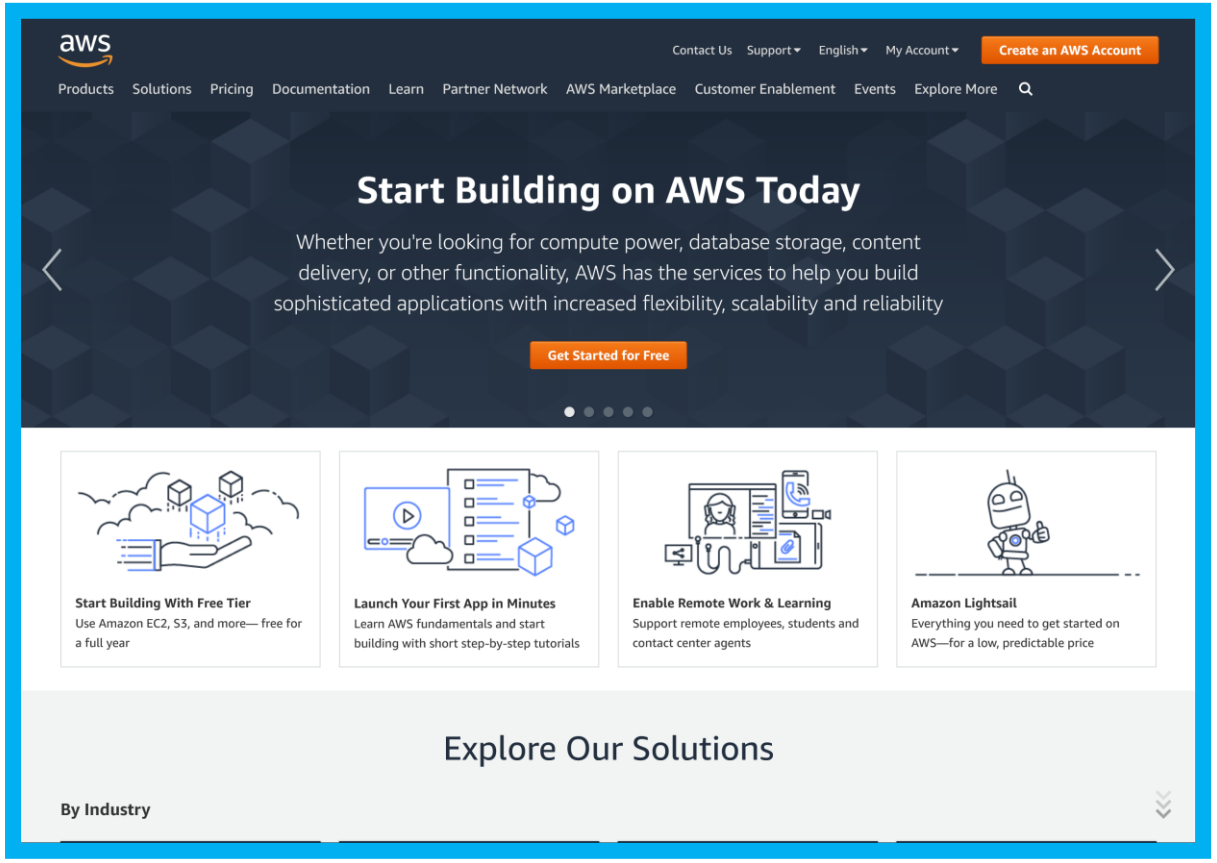


Concept of *almost-anything-as-a-service*

Delivers computing, storage, apps, etc. as a service across a network

Can be thought as an extension of virtualization as they use virtualization as base for their functionality

- Amazon EC2 has thousands of servers, **millions of virtual machines**, petabytes of storage available across the Internet, where users pay based on usage

# Cloud Computing

**Cloud computing environments composed of traditional operating systems + VMMs + cloud management tools**

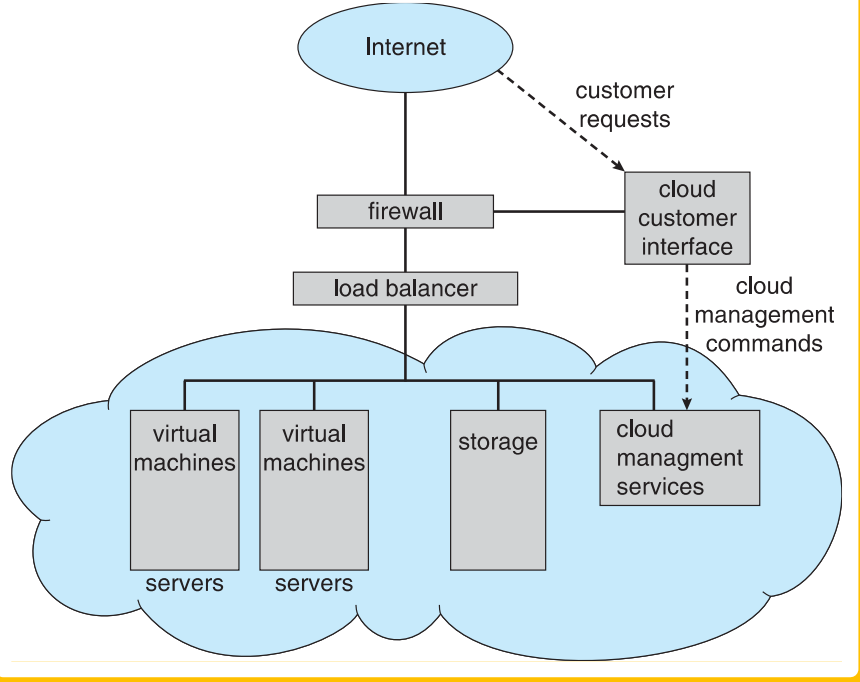

# Cloud Computing

## Almost-anything-as-a-service

**Software-as-a-Service (SaaS):** One or more applications available via the Internet (e.g., Google Docs, Dropbox, etc.)

**Platform-as-a-Service (PaaS):** Software stack ready for application use via the Internet (e.g., AWS Elastic Beanstalk, Google App Engine, etc.)

**Infrastructure-as-a-Service (IaaS):** Servers or storage available over Internet (e.g., AWS EC2, DigitalOcean, etc.)

## Cloud types

**Public cloud:** Available via Internet to anyone willing to pay

**Private cloud:** Run by a company for the company's own use

**Hybrid cloud:** Includes both public and private cloud components

# Embedded Systems

**Most prevalent** form of computers

- Vary considerably depending on purpose
- Operating system may be limited-purpose, etc.

Some systems may have operating system, some perform tasks without

**Real-time operating system** has well-defined fixed time constraints

- Processing must be done within constraint
- Correct operation only if constraints met

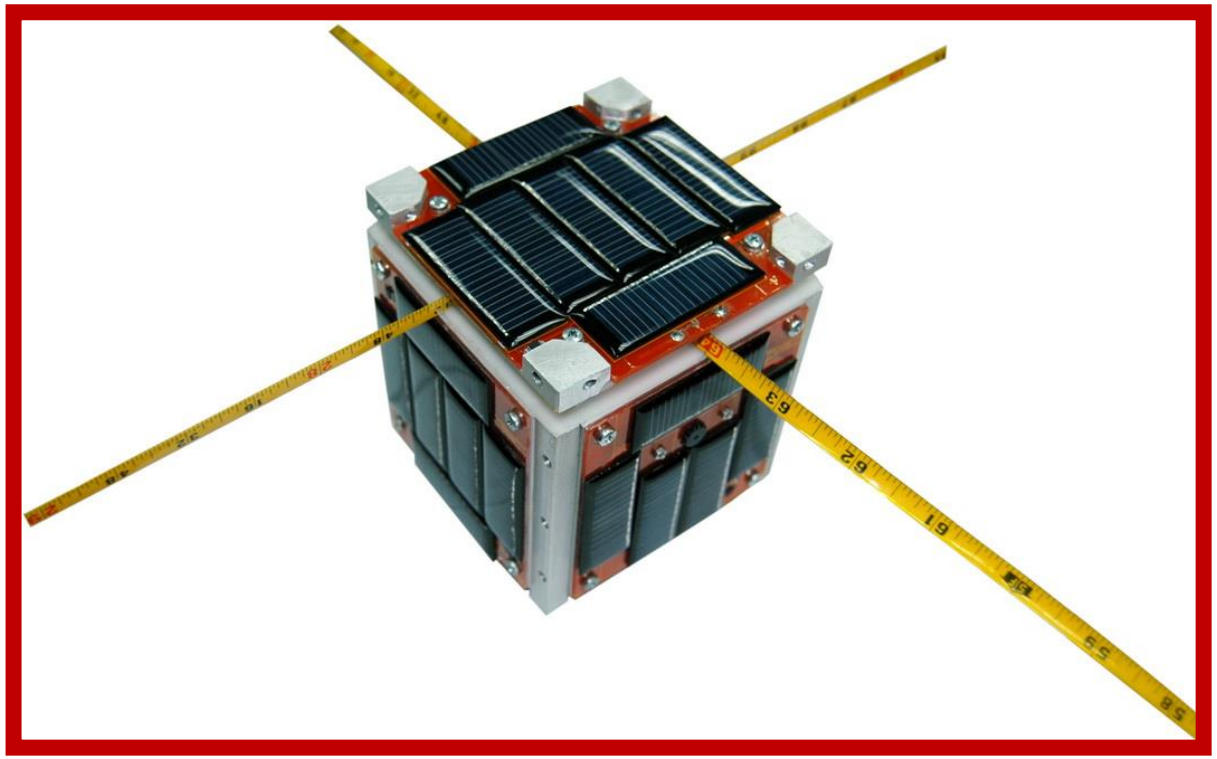

# Questions?
# Thank You!


Weihan.Goh {at} Singaporetech.edu.sg

https://www.singaporetech.edu.sg/directory/faculty/weihan-goh

https://sg.linkedin.com/in/weihan-goh