

# Processes

## CSD2180 Operating Systems

BSc in Computer Science (IMGD / RTIS)

Singapore Institute of Technology / DigiPen Institute of Technology  
September 2021

# Attendance Taking

<https://forms.office.com/r/rC3TKbsgpR>

- Log in to your SIT account to submit the form
- You can only submit the form once
- The codeword is **tungsten**



## Week 2 Quiz Details

**Open book**

**No need to use the Respondus Lockdown Browser**



# Process Concept



A process is a program  
in execution 



# The Process

An operating system executes a variety of programs that run as a process

Program is a **passive** entity stored on disk (**executable file**); process is **active**

Program becomes process when an executable file is loaded into memory

## Process


- A program in execution
- Process execution must progress in sequential fashion
- No parallel execution of instructions of a single process

## One program can be several processes

- Consider multiple users executing the same program



# Process in Memory




**Text section:** Comprises the compiled program code (from non-volatile storage) when launched

**Stack:** Contains temporary data such as function parameters, return addresses, local variables

**Data section:** Contains global and static variables, allocated and initialized prior to executing main

**Heap:** Contains memory dynamically allocated during run time and is managed via calls such as malloc, free, etc.

# Process in Memory



**Stack and heap grow toward one another; the operating system must ensure they do not overlap one another**

## Should they ever meet

- Stack overflow error will occur
- A call to new or malloc will fail due to insufficient memory available

## Some notes about process in memory

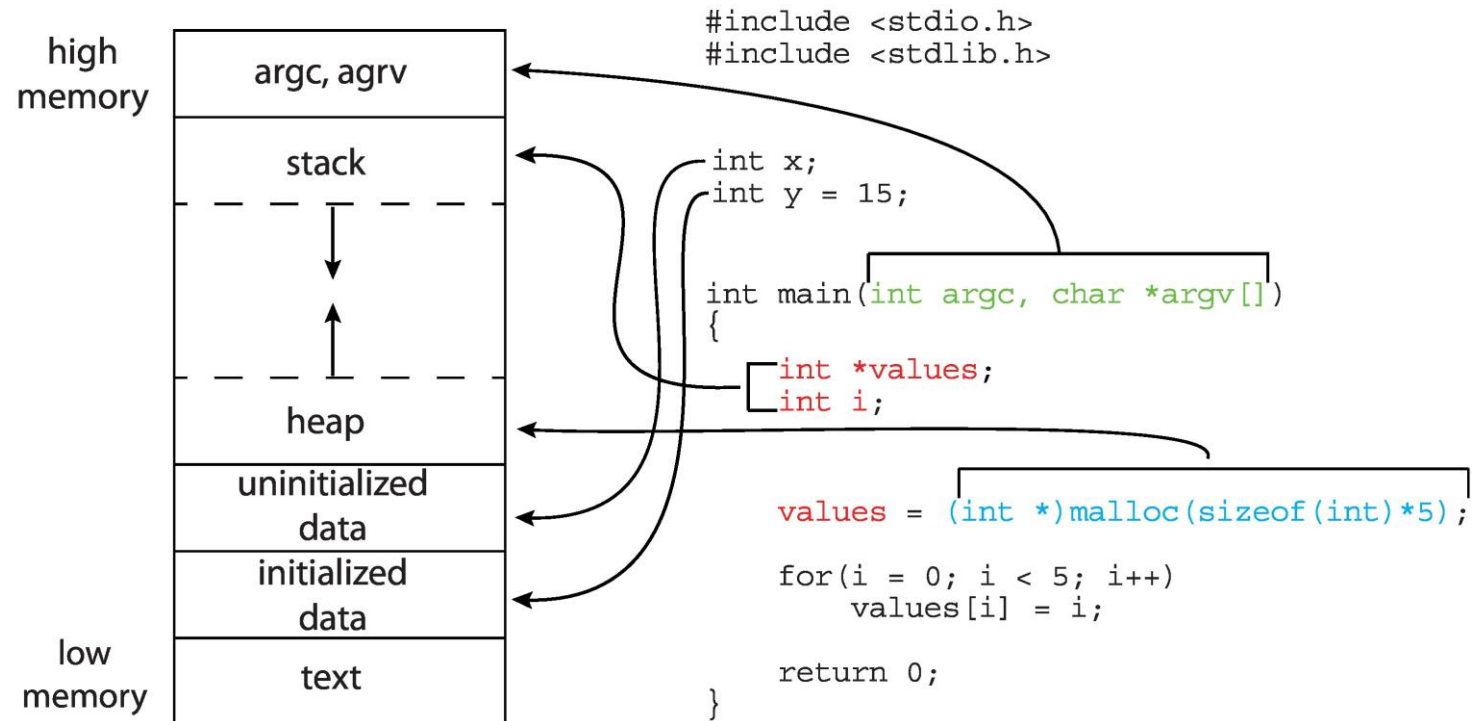
- Sizes of text and data sections do not change during program run time
- Stack and heap sections can shrink and grow dynamically during execution
- Each time a function is called, an **activation record** containing function parameters, local variables, and return address is pushed onto the stack
- When control is returned from the function, the activation record is popped from the stack
- Heap will grow as memory is dynamically allocated, and will shrink when memory is returned to the system



# Memory Layout of a C Program

Similar to the previous diagram, except

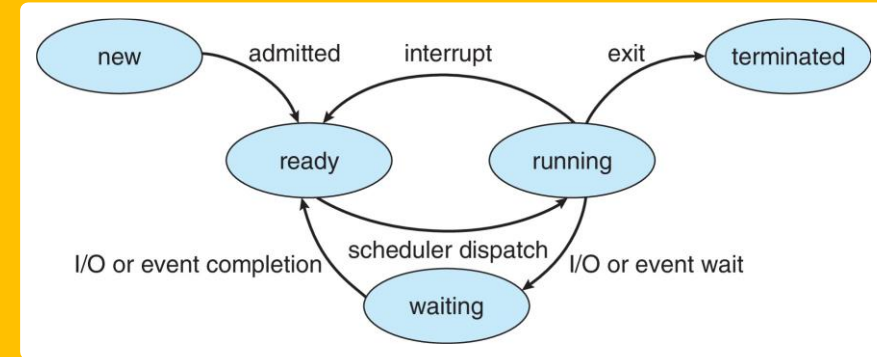
- The data section is divided into distinct sections for initialized and uninitialized data
- A separate section is provided for the *argc* and *argv* parameters passed to the *main()* function



# Process State

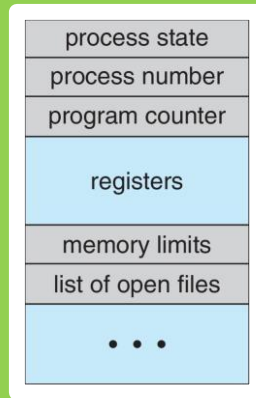
## A process changes state as it executes

- **New:** The process is being created
- **Ready:** The process is waiting to be assigned to a processor
- **Running:** Instructions are being executed
- **Waiting:** The process is waiting for some event to occur
- **Terminated:** The process has finished execution



**Protip:** Many processes may be ready and waiting, but only one process can be running on any CPU core at any instant

# Process Control Block




PCB serves as the repository for all data needed to start, or restart, a process, along with some accounting data

## Information associated with each process

- **Process state:** Running, waiting, etc.
- **Process number:** Process ID and parent process ID
- **Program counter:** Location of instruction to next execute
- **CPU registers:** Contents of all process centric registers
- **CPU scheduling information:** Priorities, scheduling queue pointers
- **Memory management information:** Memory allocated to the process
- **Accounting information:** CPU used, clock time elapsed since start, time limits
- **I/O status information:** I/O devices allocated to process, list of open files



# Process Scheduling

When can the operating  
system switch the CPU  
from one process to  
another? 



Which process should it  
switch to? 



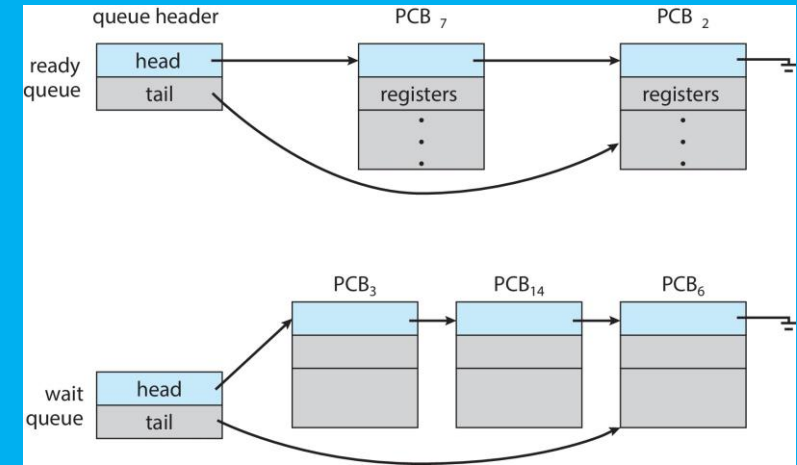
# Process Scheduling

**Process scheduler** selects among available processes for next execution on CPU core

Maintains scheduling queues of processes

- **Ready queue:** Set of all processes residing in main memory, ready and waiting to execute
- **Wait queues:** Set of processes waiting for an event (i.e., I/O)

Processes migrate among the various queues



**Goal: Maximize CPU use, quickly switch (ready) processes onto CPU core**



# Process Scheduling

**Goal: Maximize CPU use, quickly switch (ready) processes onto CPU core**

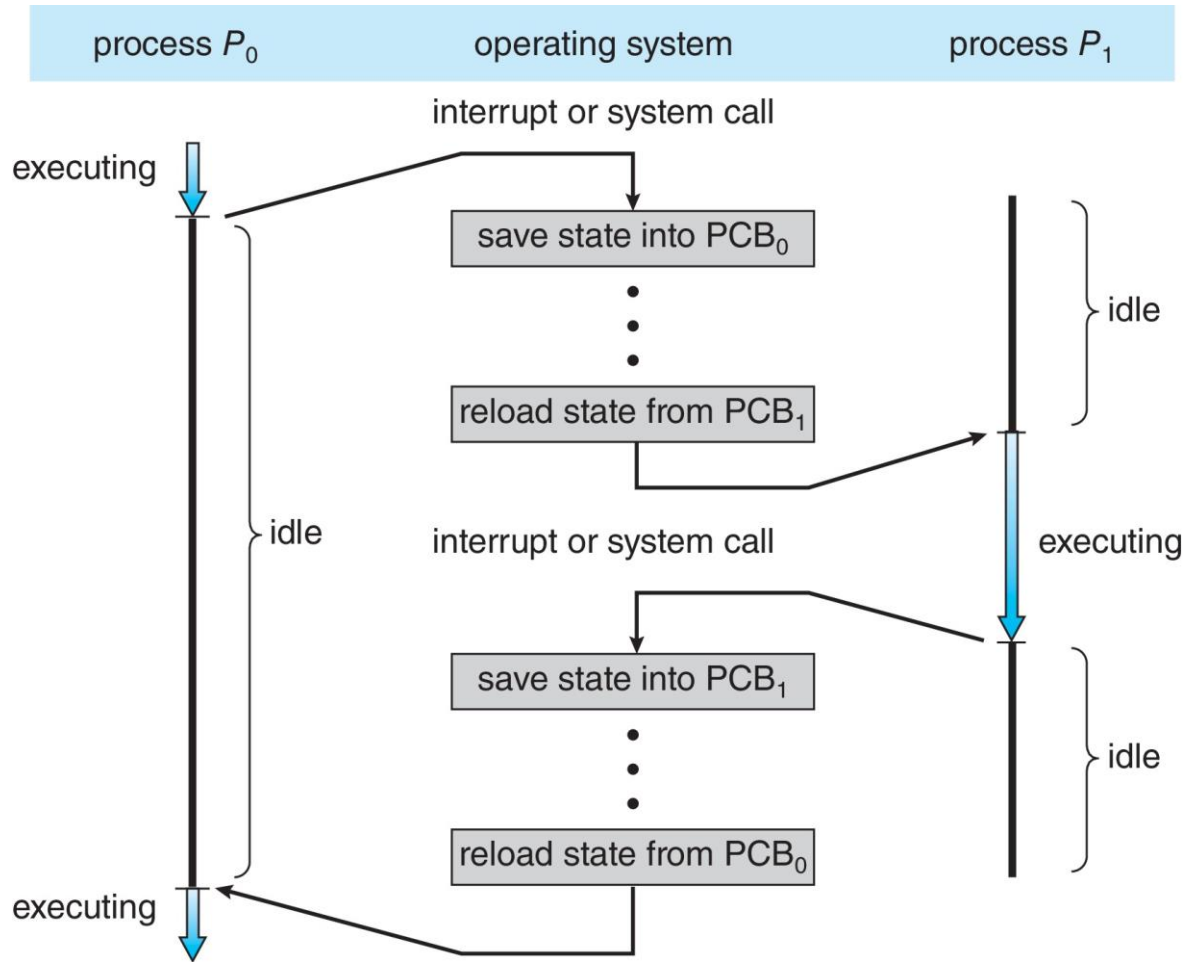
Processes can be described as either

- **I/O bound:** Process spends more time doing I/O than computations; many short CPU bursts
- **CPU bound:** Process spends more time doing computations; few very long CPU bursts

## Constraints and considerations

- There is a need to deliver "acceptable" response times for all programs, particularly interactive ones
- Process scheduler must implement suitable policies for swapping processes in and out of the CPU
- Every time the operating system steps in to swap processes, it takes up CPU time which is "lost" for doing any productive work

# Context Switch



When CPU switches to another process, the system must **save the context** of the old process and **load the saved context** for the new process via a context switch

Context of a process represented in the PCB

Context switching time is **pure overhead**; the system does no useful work while switching

The more complex the operating system and the PCB, the longer the context switch takes



# Operations on Processes

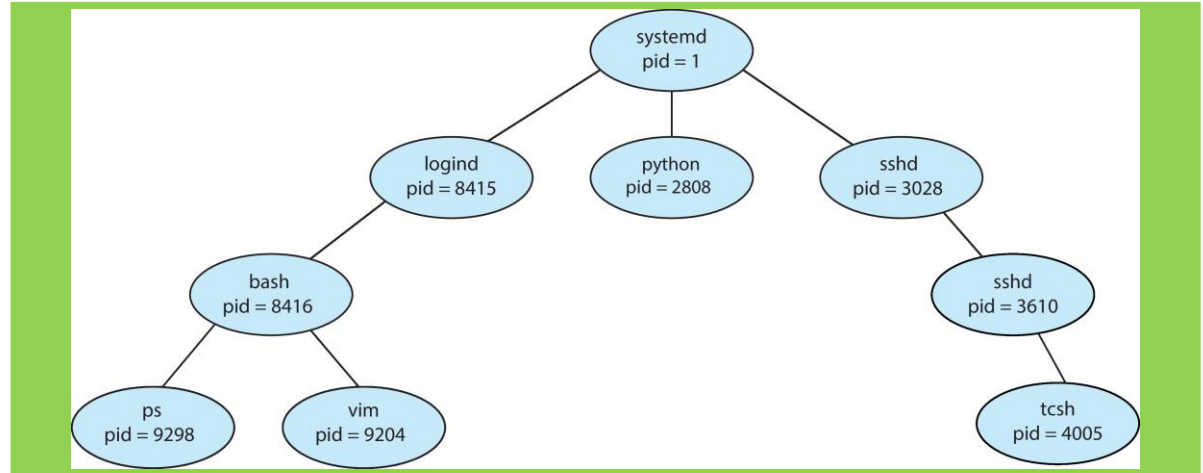
Operating system must  
provide mechanisms for  
process creation and  
process termination

# Process Creation

A process may create new processes over its course of execution

The creating process is called the **parent** process; the new processes are called the **child** of that process

Each of these new processes may in turn create other processes, forming a tree of processes



Processes are generally identified and managed via a *process identifier* (PID)

# Process Creation

## Resource sharing options

- Parent and children share all resources
- Children share subset of parent's resources
- Parent and child share no resources

## Execution options


- Parent and children execute concurrently
- Parent waits until children terminate

## Address space options

- Child duplicate of parent
- Child has a program loaded into it

# Creating a Process in Unix

- **fork()** system call creates new process
- **exec()** system call used after a **fork()** to replace the process' memory space with a new program
- Parent process calls **wait()** waiting for the child to terminate



```

#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    /* return 0 to child process, non-zero child PID to the parent */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    } else if (pid == 0) { /* child process */
        /* execlp() is a version of the exec() syscall */
        execlp("/bin/ls", "ls", NULL);
    } else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete"); /* any last words? */
    }

    return 0;
}

```



# Process Termination

A process terminates when it finishes executing its **final statement** and asks the operating system to delete it by using the **exit()** system call

The process may return a **status value** to its waiting parent process (via **wait()**)

Process resources, including physical and virtual memory, open files, and I/O buffers, are deallocated and reclaimed by the operating system

```

weihan@DESKTOP-7AVEQ08: /r  +  v
(Message from Kali developers)
We have kept /usr/bin/python pointing to Python 2 for backwards
compatibility. Learn how to change this and avoid this message:
= https://www.kali.org/docs/general-use/python3-transition/
(Run: "touch ~/.hushlogin" to hide this message)
(weihan@DESKTOP-7AVEQ08)-[/mnt/c/Users/Goh Weihan]
$ cd Desktop

(weihan@DESKTOP-7AVEQ08)-[/mnt/c/Users/Goh Weihan/Desktop]
$ gcc -o my_app my_app.c

(weihan@DESKTOP-7AVEQ08)-[/mnt/c/Users/Goh Weihan/Desktop]
$ ./my_app

(weihan@DESKTOP-7AVEQ08)-[/mnt/c/Users/Goh Weihan/Desktop]
$ echo $?
0

(weihan@DESKTOP-7AVEQ08)-[/mnt/c/Users/Goh Weihan/Desktop]
$ ./my_app
^C

(weihan@DESKTOP-7AVEQ08)-[/mnt/c/Users/Goh Weihan/Desktop]
$ echo $?
130

(weihan@DESKTOP-7AVEQ08)-[/mnt/c/Users/Goh Weihan/Desktop]
$

```

# Process Termination

## Parent may terminate the execution of child processes for various reasons


- Child has **exceeded allocated resources**
- Task assigned to child is **no longer required**
- The parent is exiting, and the operating systems **does not allow a child to exist** if its parent terminates

## Some operating systems do not allow child to exists if its parent has terminated

- If a process terminates, then all its children must also be terminated
- **Cascading termination:** All child processes, grandchild processes, etc., are terminated
- Such termination is initiated by the operating system

# Process Termination

A zombie process is one that has terminated, but whose parent has yet to invoke `wait()`



Orphaned processes are child processes of a parent that terminated without invoking `wait()`



# What happens to orphaned processes?



# Interprocess Communication

# Interprocess Communication

Processes within a system may be independent or cooperating

- **Independent processes** cannot affect or be affected by the execution of another process
- **Cooperating processes** can affect or be affected by other processes, including sharing data

## Two models for interprocess communication

- **Shared memory**
- **Message passing**

## Reasons for cooperating processes


- **Information sharing**
- **Computation speedup**
- **Modularity**
- **Convenience**

**Cooperating processes need interprocess communication (IPC)**

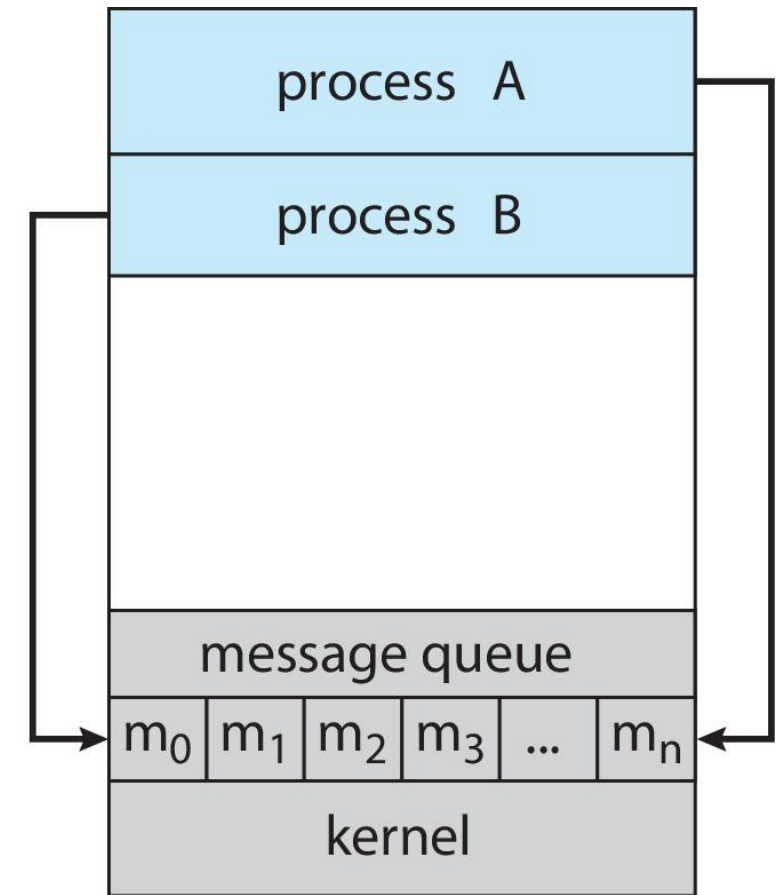
# Interprocess Communication Models

## Two models for interprocess communication

- Shared memory (a), where a region of memory is shared by the cooperating processes
- Message passing (b), where messages exchanges happen between the cooperating processes




(a)



(b)



# Shared Memory



An area of memory shared among processes that wish to communicate

Communication is controlled by the processes, not the operating system

What can possibly go  
wrong? 🤔

# What Can Possibly Go Wrong with Shared Memory?

**A process tries to write to an already full shared memory buffer**

**A process tries to read from an empty shared memory buffer**

**Multiple processes try to access the shared memory buffer at the same time**



# How do we solve this?





# Let's take a step back...



# Producer-Consumer Problem

A model for cooperating processes



**Producer process** produces information



**Consumer process** consumes information

Two variations of the producer-consumer problem

- **Unbounded-buffer**, i.e., no limit on the buffer size
  - Producer never waits
  - Consumer waits if there is nothing in buffer to consume
- **Bounded-buffer**, i.e., fixed buffer size
  - Producer must wait if buffer is full
  - Consumer waits if there is nothing in buffer to consume

# Message Passing

Processes communicate with each other by passing messages to each other

IPC facility provides two operations

- **send(message)**
- **receive(message)**

If two processes wish to communicate, they need to

- Establish a **communication link** between them
- Exchange messages via send / receive

## Implementation considerations

- Direct or indirect communication?
- Synchronous or asynchronous communication?
- Automatic or explicit buffering?



# Direct and Indirect Communication in Message Passing



## Direct communication

- Processes must name each other explicitly
  - **send(P, message)**: Send message to process P
  - **receive(Q, message)**: Receive message from process Q
- Properties of direct communication link
  - Links are established automatically
  - A link is associated with exactly one pair of communicating processes
  - Between each pair there exists exactly one link
  - The link may be unidirectional, but is usually bi-directional



## Indirect communication

- Messages are directed and received from *mailboxes* (also called *ports*)
  - Each mailbox has a **unique id**
  - Processes can communicate only if they share a mailbox
- Properties of indirect communication link
  - Link established only if processes share a common mailbox
  - A link may be associated with many processes
  - Each pair of processes may share several communication links
  - Link may be unidirectional or bi-directional

# A Note on Indirect Communication in Message Passing

## Operations

- Create a new mailbox (port)
- Send and receive messages through mailbox
- Delete a mailbox

## How to communicate?

- **send(Z, message)**: Send message to mailbox Z
- **receive(Z, message)**: Receive message from mailbox Z

## Mailbox sharing

- Problem
  - P, Q, and R share mailbox A
  - P sends; Q and R receive
  - Who gets the message?
- Possible solutions
  - Allow a link to be associated with at most two processes
  - Allow only one process at a time to execute a receive operation
  - Allow the system to select arbitrarily the receiver; sender is then notified who the receiver was

# Synchronization

Message passing may be either blocking or non-blocking

Different combinations possible

If both send and receive are **blocking**, we have a **rendezvous**\* between sender and receiver

\* **Rendezvous**: In interprocess communication, when blocking mode is used, the meeting point at which a send is picked up by a receive

## Blocking is considered synchronous

- **Blocking send**: The sender is blocked until the message is received
- **Blocking receive**: The receiver is blocked until a message is available

## Non-blocking is considered asynchronous

- **Non-blocking send**: the sender sends the message and continue
- **Non-blocking receive**: the receiver receives either a valid message, or null message

# Buffering

Messages exchanged by communicating processes reside in a temporary queue

Such queue can be implemented in three ways

- **Zero capacity:** Queue has a maximum length of zero; sender must wait for receiver
- **Bounded capacity:** Queue has finite length  $n$ ; sender must wait if link is full
- **Unbounded capacity:** Queue length is (potentially) infinite; sender never waits


# Questions? Thank You!

 [Weihan.Goh {at} Singaporetech.edu.sg](mailto:Weihan.Goh@Singaporetech.edu.sg)

 <https://www.singaporetech.edu.sg/directory/faculty/weihan-goh>

 <https://sg.linkedin.com/in/weihan-goh>

```
define traverseLinkedList(headPointer):
    myID = "11111111111111111111"
    authToken = "11111111111111111111"
    museumAddress = "11111111111111111111"
    client = mailRestClient(myID, authToken)
    client.messages.send(to=museumAddress,
        subj="Item donation?", body="Thought you
        might be interested: "+str(headPointer))
    return
```



CODING INTERVIEW TIP: INTERVIEWERS GET REALLY MAD WHEN YOU TRY TO DONATE THEIR LINKED LISTS TO A TECHNOLOGY MUSEUM.