

# Deadlocks, Part II

**CSD2180 Operating Systems**

BSc in Computer Science (IMGD / RTIS)

Singapore Institute of Technology / DigiPen Institute of Technology

September 2021

# Attendance Taking

<https://forms.office.com/r/UfBi6v7u8E>

- Log in to your SIT account to submit the form
- You can only submit the form once
- The codeword is **tablet**



Several small, colorful geometric shapes are scattered around the slide: a yellow circle in the upper left, a purple square in the upper center, a blue triangle in the upper right, a blue triangle in the lower left, and a yellow circle in the lower right.

# Deadlock Avoidance

# Deadlock Avoidance

Requires that the system has some additional *a priori* information available


- Simplest and most useful model requires that each process declare the **maximum number** of resources of each type that it may need
- The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition
- Resource-allocation state is defined by the number of available and allocated resources, and the maximum demands of the processes

## Safe State

- When a process requests an available resource, system must decide if immediate allocation leaves the system in a *safe state*
- System is in safe state if there exists a sequence  $\langle P_1, P_2, \dots, P_n \rangle$  of all processes in the systems such that for each  $P_i$ , the resources that  $P_i$  can still request can be satisfied by currently available resources + resources held by all  $P_j$ , where  $j < i$
- That is
  - If  $P_i$  resource needs are not immediately available, then  $P_i$  can wait until all  $P_j$  have finished
  - When  $P_j$  is finished,  $P_i$  can obtain needed resources, execute, return allocated resources, and terminate
  - When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on



# Safe State and Deadlocks



If a system is in safe state, it has no deadlocks

If a system is in unsafe state, there is possibility of deadlock

Deadlock avoidance algorithms ensure that a system will never enter an unsafe state

# Safe State and Deadlocks

Consider a system with twelve units of resources and three threads

Is this system in a safe state?

Thread	Maximum Resource Needs	Current Resource Needs
$T_A$	10	5
$T_B$	4	2
$T_C$	9	2

# Safe State and Deadlocks

Consider a system with twelve units of resources and three threads

Is this system in a safe state?

Thread	Maximum Resource Needs	Current Resource Needs
$T_A$	10	5
$T_B$	4	2
$T_C$	9	2

## Situation analysis

- Nine units of resources currently used; 3 units free
- $T_B$  can be allocated all its remaining resources (2 units) and then returns them all when done
  - 5 units will be available when  $T_B$  is done
- Then,  $T_A$  can be allocated all its remaining resources (5 units) and return them all when done
  - 10 units will be available when  $T_A$  is done
- Finally,  $T_C$  can be allocated all its remaining resources (7 units) and return them all when done
  - 12 units will be available when done

The sequence  $\langle T_B, T_A, T_C \rangle$  satisfies the safety condition

# Safe State and Deadlocks

Consider another system with twelve units of resources and three threads

Is this system in a safe state?

Thread	Maximum Resource Needs	Current Resource Needs
$T_A$	10	5
$T_B$	4	2
$T_C$	9	3



# Safe State and Deadlocks

Consider another system with twelve units of resources and three threads

Is this system in a safe state?

Thread	Maximum Resource Needs	Current Resource Needs
$T_A$	10	5
$T_B$	4	2
$T_C$	9	3

## Situation analysis


- Ten units of resources currently used; 2 units free
- $T_B$  can be allocated all its remaining resources (2 units) and then returns them all when done
  - 4 units will be available when  $T_B$  is done
- No other threads can obtain all its resources
  - $T_A$  requires 5 units more, while  $T_C$  requires 6 units more

**The system is not in a safe state!**

# Deadlock Avoidance Algorithms

If only single instances of each resource type, use a *resource-allocation graph*

If multiple instances of a resource type exist, use the *banker's algorithm*



# Resource-Allocation Graph Approach

A **claim edge**  $P_i \rightarrow R_j$  indicates that process  $P_j$  may request resource  $R_j$


- This is represented by a dashed line

## Edge characteristics


- Claim edge converts to request edge when a process requests a resource
- Request edge converted to an assignment edge when the resource is allocated to the process
- When a resource is released by a process, assignment edge reconverts to a claim edge

Resources must be claimed *a priori* in the system

- Before process starts executing, all its claim edges must already appear in the resource-allocation graph




# Resource-Allocation Graph Approach



Request can be granted only if converting the request edge to an assignment edge does not result in formation of a cycle

Should  $T_2$  be assigned  $R_2$  if it requests for it?

# Resource-Allocation Graph Approach



Request can be granted only if converting the request edge to an assignment edge does not result in formation of a cycle

Should  $T_2$  be assigned  $R_2$  if it requests for it?

Not a good idea since a cycle will form, indicating that a deadlock will occur...

# Banker's Algorithm

Multiple instances of resources

Each process must *a priori* claim maximum use

When a process requests a resource, it may have to wait

When a process gets all its resources it must return them in a finite amount of time

Think of it from the perspective of a bank lending money...



# Banker's Algorithm

## Data structures for the banker's algorithm

- Let
  - $n$  = number of processes
  - $m$  = number of resources types
- **Available:** Vector of length  $m$  indicating the number of available resources of each type
  - If  $\text{Available}[j] = k$ , there are  $k$  instances of resource type  $R_j$  available
- **Max:**  $n \times m$  matrix ( $n$  rows,  $m$  columns) defining the maximum demand of each process
  - If  $\text{Max}[i, j] = k$ , then process  $P_i$  will request at most  $k$  instances of resource type  $R_j$
- **Allocation:**  $n \times m$  matrix defining the number of resources of each type currently allocated to each process
  - If  $\text{Allocation}[i, j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$
- **Need:**  $n \times m$  matrix indicating the remaining resource need of each process
  - If  $\text{Need}[i, j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$
- It can be seen that  $\text{Need}[i, j] = \text{Max}[i, j] - \text{Allocation}[i, j]$

# Banker's Algorithm

## Notation for banker's algorithm

- Let  $X$  and  $Y$  be vectors of length  $n$ 
  - We say that  $X \leq Y$  if and only if  $\forall i = 1, 2, \dots, n, X[i] \leq Y[i]$
  - E.g., If  $S = (1, 7, 3, 2)$  and  $T = (0, 3, 2, 1)$ , then  $T \leq S$
- Let each row in the matrices **Allocation** and **Need** be vectors, and refer to them as  $\text{Allocation}_i$  and  $\text{Need}_i$

# Safety Algorithm (Banker's Algorithm)

Let **Work** and **Finish** be vectors of length  $m$  and  $n$ , respectively

## 1. Initialize

$\text{Work} = \text{Available}$   
 $\text{Finish}[i] = \text{false}, \quad \forall i = 1, 2, \dots, n$

## 2. Find an $i$ such that

$(\text{Finish}[i] = \text{false}) \wedge (\text{Need}_i \leq \text{Work})$

- If no such  $i$  exists, go to **Step 4**

## 3. Run the following, then go to **Step 2**

$\text{Work} = \text{Work} + \text{Allocation}_i$   
 $\text{Finish}[i] = \text{true}$

## 4. If $\forall i, \text{Finish}[i] == \text{true}$ , then the system is in a safe state

# Resource-Request Algorithm for Process $P_i$ (Banker's Algorithm)

$\text{Request}_i$  = request vector for process  $P_i$

- If  $\text{Request}_i[j] = k$ , then process  $P_i$  wants  $k$  instances of resource type  $R_j$

1. If  $\text{Request}_i \leq \text{Need}_i$  go to **Step 2**
  - Otherwise, raise error condition, since process has exceeded its maximum claim
2. If  $\text{Request}_i \leq \text{Available}$ , go to **Step 3**
  - Otherwise  $P_i$  must wait, since resources are not available

3. Pretend to allocate requested resources to  $P_i$  by modifying the state as follows

$$\begin{aligned}\text{Available} &= \text{Available} - \text{Request}_i \\ \text{Allocation}_i &= \text{Allocation}_i + \text{Request}_i \\ \text{Need}_i &= \text{Need}_i - \text{Request}_i\end{aligned}$$

4. Check for safety
  - If safe, the resources are allocated to  $P_i$
  - If unsafe,  $P_i$  must wait, and the old resource-allocation state is restored

# Banker's Algorithm Example

5 processes,  $P_0$  through  $P_4$

3 resources

- 10 instances of resource A
- 5 instances of resource B
- 7 instances of resource C
- **Available** = (3, 3, 2)

	Allocation			Max			Need		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	7	5	3			
$P_1$	2	0	0	3	2	2			
$P_2$	3	0	2	9	0	2			
$P_3$	2	1	1	2	2	2			
$P_4$	0	0	2	4	3	3			

# Banker's Algorithm Example

5 processes,  $P_0$  through  $P_4$

3 resources

- 10 instances of resource A
- 5 instances of resource B
- 7 instances of resource C
- **Available** = (3, 3, 2)

	Allocation			Max			Need		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	7	5	3	7	4	3
$P_1$	2	0	0	3	2	2	1	2	2
$P_2$	3	0	2	9	0	2	6	0	0
$P_3$	2	1	1	2	2	2	0	1	1
$P_4$	0	0	2	4	3	3	4	3	1



# Banker's Algorithm Example

5 processes,  $P_0$  through  $P_4$

3 resources

- 10 instances of resource A
- 5 instances of resource B
- 7 instances of resource C
- **Available** = (3, 3, 2)

	Allocation			Max			Need		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	7	5	3	7	4	3
$P_1$	2	0	0	3	2	2	1	2	2
$P_2$	3	0	2	9	0	2	6	0	0
$P_3$	2	1	1	2	2	2	0	1	1
$P_4$	0	0	2	4	3	3	4	3	1

**Available** = **Work** = (3, 3, 2)

$P_1: (1, 2, 2) \leq (3, 3, 2)$   
 ○ When done, Work = (5, 3, 2)

$P_3: (0, 1, 1) \leq (5, 3, 2)$   
 ○ When done, Work = (7, 4, 3)

$P_4: (4, 3, 1) \leq (7, 4, 3)$   
 ○ When done, Work = (7, 4, 5)

$P_2: (6, 0, 0) \leq (7, 4, 5)$   
 ○ When done, Work = (10, 4, 7)

$P_0: (7, 4, 3) \leq (10, 4, 7)$   
 ○ When done, Work = (10, 5, 7)

The system is in a **safe state** since the sequence  $\langle P_1, P_3, P_4, P_2, P_0 \rangle$  satisfies safety criteria

# Banker's Algorithm Example

Suppose now at this safe state,  $P_1$  requests for  $(1, 0, 2)$

Can this be fulfilled?

	Allocation			Max			Need		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	7	5	3			
$P_1$	3	0	2	3	2	2			
$P_2$	3	0	2	9	0	2			
$P_3$	2	1	1	2	2	2			
$P_4$	0	0	2	4	3	3			

# Banker's Algorithm Example

Suppose now at this safe state,  $P_1$  requests for  $(1, 0, 2)$

## Perform checks

- $\text{Request}_1 = (1, 0, 2) \leq (1, 2, 2)$ , hence  $\text{Request}_i \leq \text{Need}_i$
- $\text{Request}_1 = (1, 0, 2) \leq (3, 3, 2)$ , hence  $\text{Request}_i \leq \text{Available}$

Pretend that request has been fulfilled to arrive at the following new state, with **Available** =  $(2, 3, 0)$

	Allocation			Max			Need		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	7	5	3			
$P_1$	3	0	2	3	2	2			
$P_2$	3	0	2	9	0	2			
$P_3$	2	1	1	2	2	2			
$P_4$	0	0	2	4	3	3			

# Banker's Algorithm Example

Suppose now at this safe state,  $P_1$  requests for  $(1, 0, 2)$

## Perform checks

- $\text{Request}_1 = (1, 0, 2) \leq (1, 2, 2)$ , hence  $\text{Request}_i \leq \text{Need}_i$
- $\text{Request}_1 = (1, 0, 2) \leq (3, 3, 2)$ , hence  $\text{Request}_i \leq \text{Available}$

Pretend that request has been fulfilled to arrive at the following new state, with **Available** =  $(2, 3, 0)$

	Allocation			Max			Need		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	7	5	3	7	4	3
$P_1$	3	0	2	3	2	2	0	2	0
$P_2$	3	0	2	9	0	2	6	0	0
$P_3$	2	1	1	2	2	2	0	1	1
$P_4$	0	0	2	4	3	3	4	3	1

# Banker's Algorithm Example

Suppose now at this safe state,  $P_1$  requests for (1, 0, 2)

## Perform checks

- $\text{Request}_1 = (1, 0, 2) \leq (1, 2, 2)$ , hence  $\text{Request}_i \leq \text{Need}_i$
- $\text{Request}_1 = (1, 0, 2) \leq (3, 3, 2)$ , hence  $\text{Request}_i \leq \text{Available}$

Pretend that request has been fulfilled to arrive at the following new state, with **Available** = (2, 3, 0)

	Allocation			Max			Need		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	7	5	3	7	4	3
$P_1$	3	0	2	3	2	2	0	2	0
$P_2$	3	0	2	9	0	2	6	0	0
$P_3$	2	1	1	2	2	2	0	1	1
$P_4$	0	0	2	4	3	3	4	3	1

**Available** = **Work** = (2, 3, 0)

$P_1$ :  $(0, 2, 0) \leq (2, 3, 0)$ ; Work = (5, 3, 2)

$P_3$ :  $(0, 1, 1) \leq (5, 3, 2)$ ; Work = (7, 4, 3)

$P_4$ :  $(4, 3, 1) \leq (7, 4, 3)$ ; Work = (7, 4, 5)

$P_0$ :  $(7, 4, 3) \leq (7, 4, 5)$ ; Work = (7, 5, 5)

$P_2$ :  $(6, 0, 0) \leq (7, 5, 5)$ ; Work = (10, 5, 7)

The system is in a **safe state** since the sequence  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  satisfies safety criteria

Continuing from here

- Can a request for (3, 3, 0) by  $P_4$  be granted?
- Can a request for (0, 2, 0) by  $P_0$  be granted?

Several small, colorful geometric shapes are scattered across the slide: a yellow circle in the upper left, a purple square in the upper center, a blue triangle in the upper right, a blue triangle in the lower left, and a yellow circle in the lower right.

# Deadlock Detection and Recovery



# Deadlock Detection and Recovery

**Goal:** Allow system to enter deadlock state

## Need to have

- Detection algorithm
- Recovery scheme

**If only single instances of a resource type, use a *wait-for graph***

**If multiple instances of a resource type exist, use a *variation of the safety algorithm* from the banker's algorithm earlier**

# Wait-For Graph

Nodes in a **wait-for graph** are processes



$P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$

- Exist if and only if corresponding resource-allocation graph contains two edges  $P_i \rightarrow R_q$  and  $R_q \rightarrow P_j$  for some resource  $R_q$

If there is a **cycle** in the wait-for graph, there exists a deadlock

**Example:** Resource-allocation graph (a) and corresponding wait-for graph (b)

- Is there a deadlock?**



# Variation of Safety Algorithm (Banker's Algorithm) for Deadlock Detection

**Available** and **Allocation** as per original algorithm

**Request:**  $n \times m$  matrix indicating the current resource request of each process

- If  $\text{Request}[i, j] = k$ , then  $P_i$  is need  $k$  more instances of  $R_j$

Let **Work** and **Finish** be vectors of length  $m$  and  $n$ , respectively

1. Initialize

$\text{Work} = \text{Available}$   
 $\text{Finish}[i] = \text{false}, \quad \forall i = 1, 2, \dots, n$

- Find an  $i$  such that  
 $(\text{Finish}[i] = \text{false}) \wedge (\text{Request}_i \leq \text{Work})$ 
  - If no such  $i$  exists, go to **Step 4**
- Run the following, then go to **Step 2**  
 $\text{Work} = \text{Work} + \text{Allocation}_i$   
 $\text{Finish}[i] = \text{true}$
- If  $\exists i, \text{Finish}[i] == \text{false}$ , then the system is in a deadlocked state
  - In particular, if  $\text{Finish}[i] == \text{false}$ , then  $P_i$  is deadlocked

# Deadlock Detection Using Variation of Safety Algorithm Example

5 processes,  $P_0$  through  $P_4$

3 resources

- 7 instances of resource A
- 2 instances of resource B
- 6 instances of resource C
- **Available** = (0, 0, 0)

	Allocation			Request		
	A	B	C	A	B	C
$P_0$	0	1	0	0	0	0
$P_1$	2	0	0	2	0	2
$P_2$	3	0	3	0	0	0
$P_3$	2	1	1	1	0	0
$P_4$	0	0	2	0	0	2

# Deadlock Detection Using Variation of Safety Algorithm Example

5 processes,  $P_0$  through  $P_4$

3 resources

- 7 instances of resource A
- 2 instances of resource B
- 6 instances of resource C
- **Available** = (0, 0, 0)

	Allocation			Request		
	A	B	C	A	B	C
$P_0$	0	1	0	0	0	0
$P_1$	2	0	0	2	0	2
$P_2$	3	0	3	0	0	0
$P_3$	2	1	1	1	0	0
$P_4$	0	0	2	0	0	2

**Available** = **Work** = (0, 0, 0)

$P_0$ : (0, 0, 0)  $\leq$  (0, 0, 0)  
 ○ When done, Work = (0, 1, 0)

$P_2$ : (0, 0, 0)  $\leq$  (0, 1, 0)  
 ○ When done, Work = (3, 1, 3)

$P_3$ : (1, 0, 0)  $\leq$  (3, 1, 3)  
 ○ When done, Work = (5, 2, 4)

$P_1$ : (2, 0, 2)  $\leq$  (5, 2, 4)  
 ○ When done, Work = (7, 2, 4)

$P_4$ : (0, 0, 2)  $\leq$  (7, 2, 4)  
 ○ When done, Work = (7, 2, 6)

The system is **not in a deadlocked state** since the sequence  $\langle P_0, P_2, P_3, P_1, P_4 \rangle$  satisfies safety criteria

# Deadlock Detection Using Variation of Safety Algorithm Example

5 processes,  $P_0$  through  $P_4$

3 resources

- 7 instances of resource A
- 2 instances of resource B
- 6 instances of resource C
- **Available** = (0, 0, 0)

	Allocation			Request		
	A	B	C	A	B	C
$P_0$	0	1	0	0	0	0
$P_1$	2	0	0	2	0	2
$P_2$	3	0	3	0	0	1
$P_3$	2	1	1	1	0	0
$P_4$	0	0	2	0	0	2



# Deadlock Detection Using Variation of Safety Algorithm Example

5 processes,  $P_0$  through  $P_4$

3 resources

- 7 instances of resource A
- 2 instances of resource B
- 6 instances of resource C
- **Available** = (0, 0, 0)

**Available** = **Work** = (0, 0, 0)

$P_0$ :  $(0, 0, 0) \leq (0, 0, 0)$ , Work = (0, 1, 0)

**Cannot continue, no suitable process available**

The system is **in a deadlocked state**

- In particular,  $P_1$ ,  $P_2$ ,  $P_3$ , and  $P_4$  are deadlocked

	Allocation			Request		
	A	B	C	A	B	C
$P_0$	0	1	0	0	0	0
$P_1$	2	0	0	2	0	2
$P_2$	3	0	3	0	0	1
$P_3$	2	1	1	1	0	0
$P_4$	0	0	2	0	0	2


# Recovery from Deadlock

## Process termination

- Abort *all* deadlocked processes?
  - Surefire way to break deadlock, but at what expense?
- Abort one process at a time until the deadlock cycle is eliminated?
  - Need to run deadlock detection after every abort
- In which order should we choose to abort?
  - Priority of the process
  - How long process has computed, and how much longer to completion
  - Resources the process has used
  - Resources process needs to complete
  - How many processes will need to be terminated
  - Is process interactive or batch?

## Resource preemption

- Selecting a victim
  - Which resources and which processes are to be preempted?
- Rollback
  - What should be done with that process?
  - Return to some safe state, then restart process from that state?
  - Restart from beginning?
- Starvation
  - How to guarantee that resources will not always be preempted from the same process?



# Questions?

# Thank You!

 [Weihan.Goh {at} Singaporetech.edu.sg](mailto:Weihan.Goh@Singaporetech.edu.sg)

 <https://www.singaporetech.edu.sg/directory/faculty/weihan-goh>

 <https://sg.linkedin.com/in/weihan-goh>

COOKING TIPS: TSP VS TBSP

TSP

TERASPOON

1,000,000,000,000  
(10<sup>12</sup>) SPOONS

TBSP

BINARY TSP

1,099,511,627,776  
(1024<sup>4</sup>) SPOONS